



# LECTURE 4:

# CONVOLUTIONAL NEURAL NETWORKS

University of Washington, Seattle


Fall 2024



# Previously in EEP 596...


Training  
Dataset

SGD




1 iteration =  
Fwd/bwd pass 1 training sample

Mini-batch



1 iteration =  
Fwd/bwd pass n-training samples  
( $n < \text{total \# of training samples}$ )



Batch GD



1 iteration =  
Fwd/bwd pass full training samples



## Previously in EEP 596...





## Previously in EEP 596...

### Optimizers

- Vanilla SGD
- Momentum
- AdaGrad
- RMSProp
- Adam



### Optimization Techniques

- Data splitting (Train/Val/Test)
- Regularization
- Data normalization
- Batch-normalization
- Network initialization
- Hyperparameter tunings



# OUTLINE

## **Part 1: Need for CNNs**

- Limitation of MLP
- Convolution Layer

## **Part 2: Convolution Filters**

- 2D convolution
- Stride
- Padding
- Volume convolutions

## **Part 3: Composing Convolutional Neural Networks**

- Convolution Layer
- Pooling Layers
- Benefits and challenges of CNNs
- Historical CNN examples

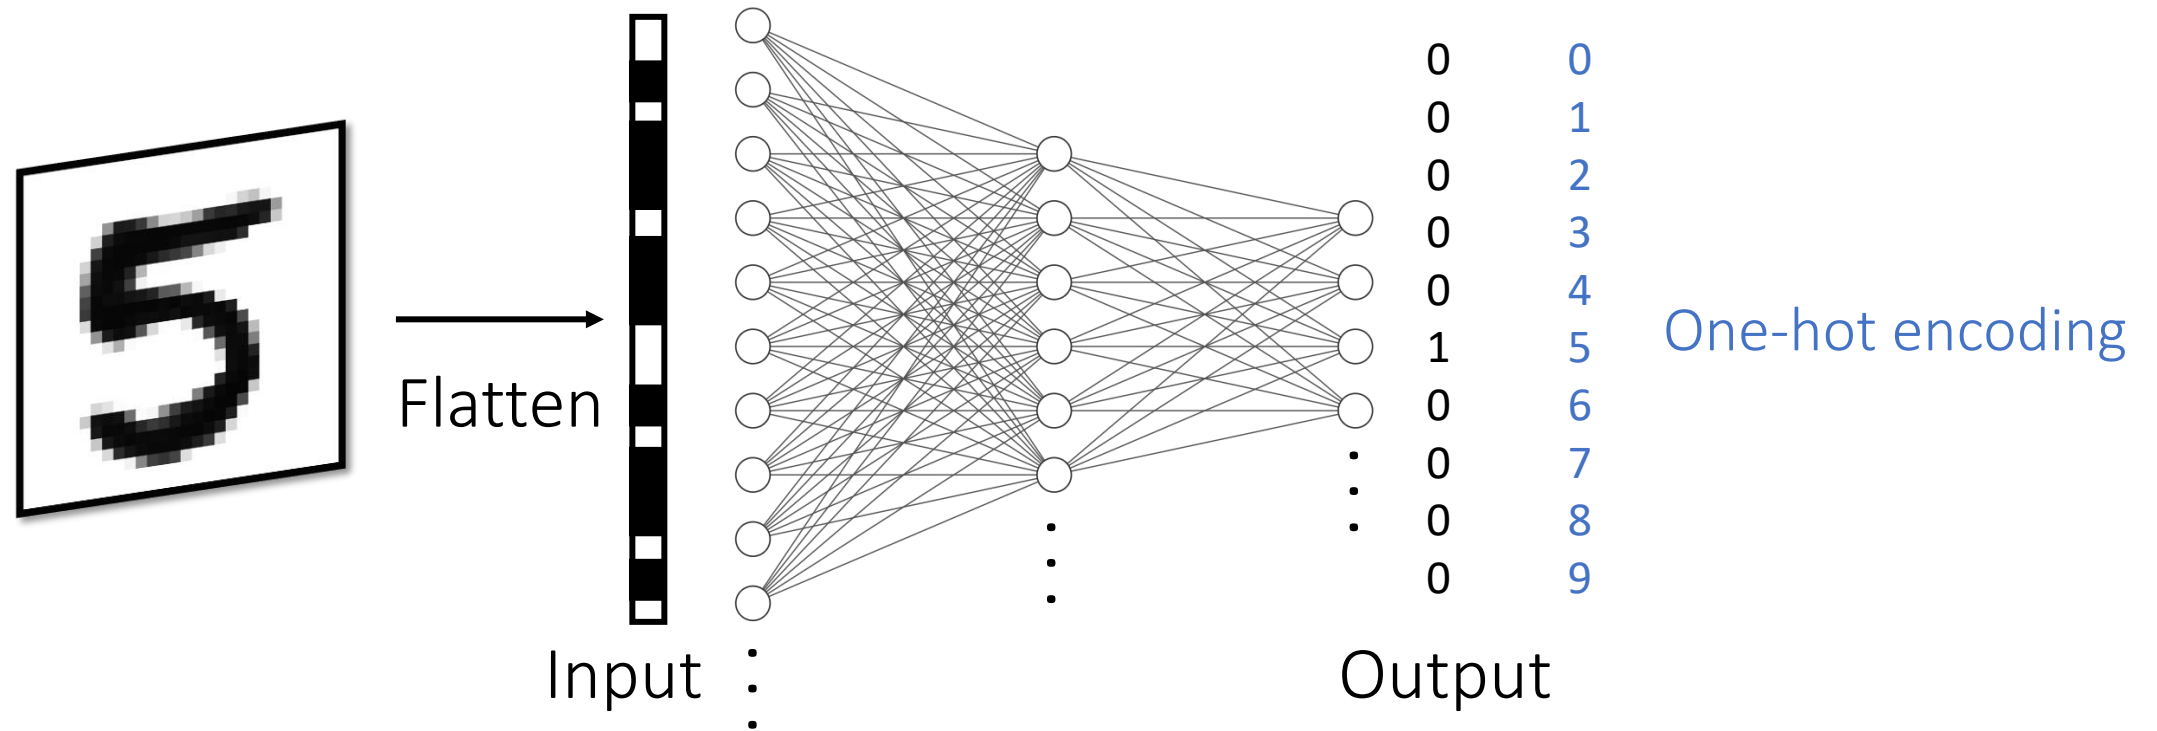


# PART 1:

## Need for CNNs

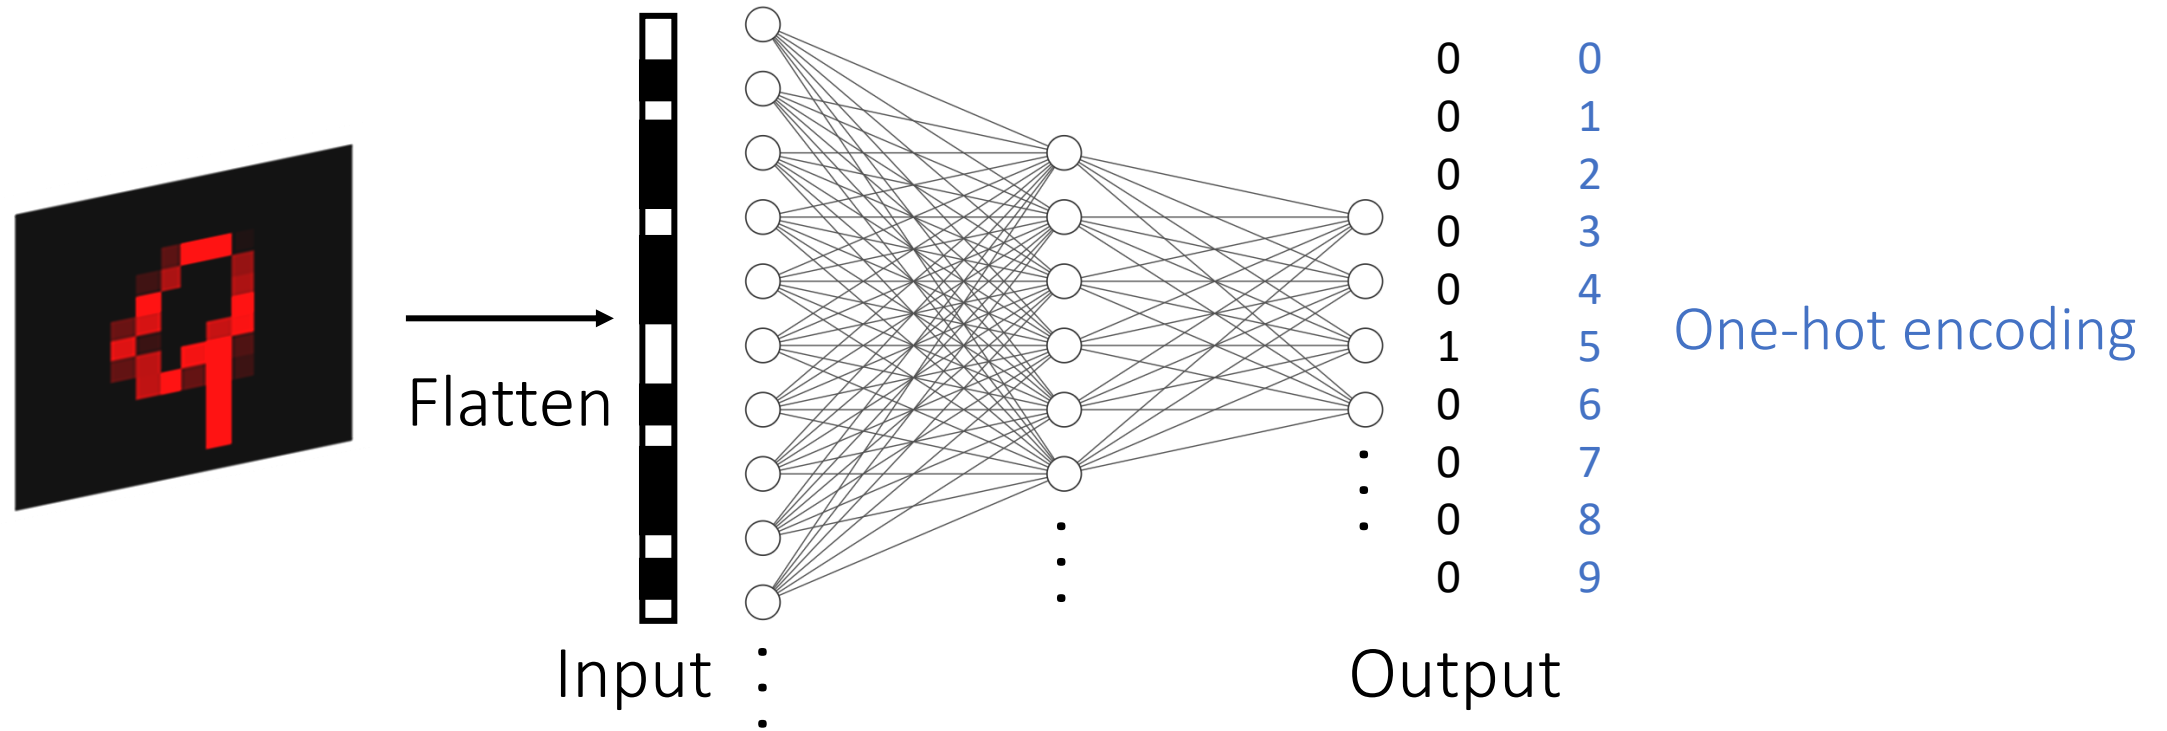


# MLP for Image Classification (Lab 3)






# MLP for Image Classification (Lab 3)








# MLP for Image Classification (Lab 3)






# MLP for Image Classification (Lab 3)



# MLP for Image Classification (Lab 3)




(3072 neurons) (4704 neurons)

$$W^{[1]} = [n^{[l-1]}, n^{[l]}] = [3072, 4704] \approx 14M$$



# MLP for Image Classification (Lab 3)



Great at Classification


Not as good with Extracting image features

Too many parameters when Flattening images

# Specialized Layers for Feature Extractions




Specialized Layers for  
Image Feature Extraction



Fully connected layers  
(Classifier)



# Full CNN Architecture



Convolution Layers + Pooling Layers  
(Image feature extraction)

Fully connected layers  
(Classifier)



# PART 2:

## Convolution Filters



# Image Convolution



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Kernel

4	3	4

Convolved Feature

$$\begin{aligned} & (1 * 1) + (0 * 0) + (0 * 1) + \\ & (0 * 1) + (1 * 1) + (0 * 0) + \\ & (1 * 1) + (0 * 0) + (1 * 1) + \end{aligned}$$












# Traditional Convolution Filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	



# Traditional Convolution Filters

CNNs **Learn** these features instead of us guessing

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	



# Traditional Convolution Filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

CNNs **Learn** these features instead of us guessing

1000 filters

$3 \times 3 = 9 \times 1000 = 9K$  parameters



# Traditional Convolution Filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

CNNs **Learn** these features instead of us guessing

1000 filters

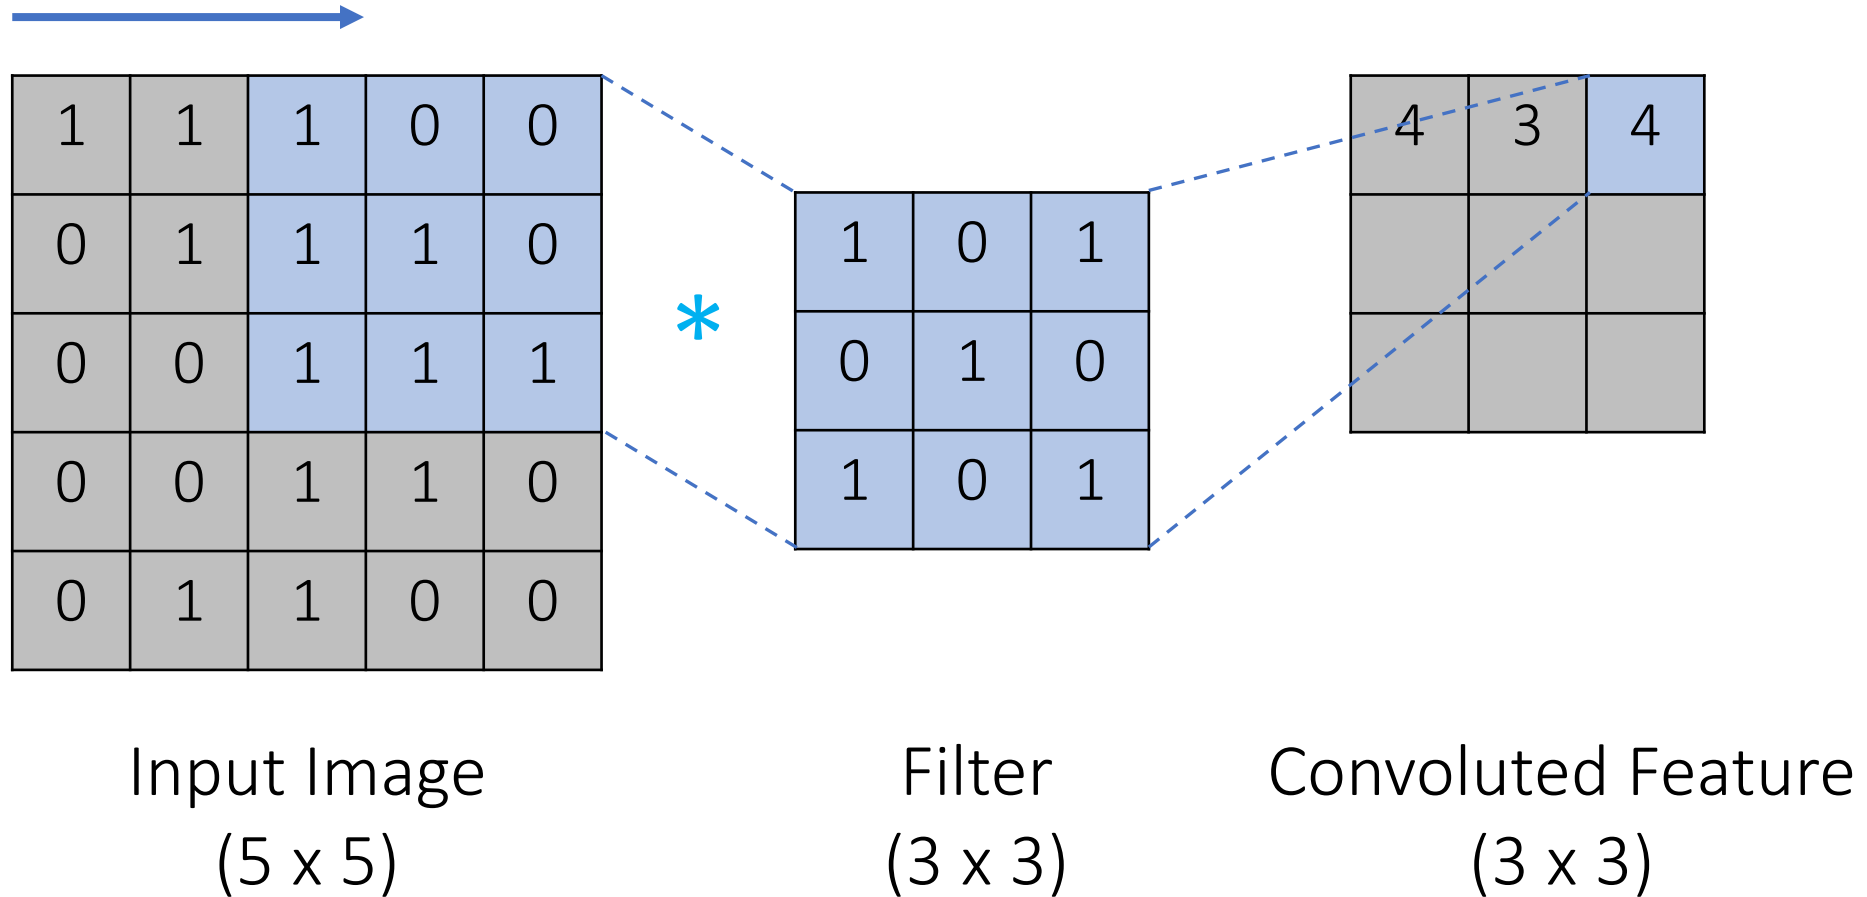
$3 \times 3 = 9 * 1000 = \mathbf{9K}$  parameters

**14M vs 9k**

Several orders of magnitude of difference in parameters




# Convolution Dimensions





# Convolution Dimensions





# Stride

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Filter


4		

Convolved Feature

Stride = 1



# Stride



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Filter

4	3	


Convolved Feature

Stride = 1





# Stride



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Filter

4	3	4


Convolved Feature

Stride = 1



# Stride

Input = 5 x 5



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Filter

Output = 3 x 3

4	3	4
3		

Convolved Feature

Stride = 1



# Stride

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Filter


4	

Convolved Feature

Stride = 2



# Stride



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Filter

4	4


Convolved Feature

Stride = 2



# Stride

Input = 5 x 5



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image

\*

1	0	1
0	1	0
1	0	1

Filter

Output = 2 x 2

4	4
2	

Convolved Feature

Stride = 2



# Padding

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image  
(5x5)



0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

Padding = 1

Padded Image  
(7x7)



# Padding

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input Image  
(5x5)




0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	1	1	1	0	0	0	0
0	0	0	1	1	1	0	0	0
0	0	0	0	1	1	1	0	0
0	0	0	0	1	1	0	0	0
0	0	0	1	1	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Padding = 2

Padded Image  
(9x9)



# Generalized Dimensions



$$(n) * (n)$$



$$\left( \frac{n + 2p - f}{s} + 1 \right) * \left( \frac{n + 2p - f}{s} + 1 \right)$$

$n$ : original image dimensions

$p$ : padding size


$f$ : filter dimension

$s$ : stride






# Volume Convolution



Input  
(5x5x3)



# Volume Convolution




Input  
(5x5x3)

Filters  
(2x2x3)

Output  
(4x4x1)

(Height x Width x Channels)

# Volume Convolution (multiple filters)






# PART 3:

## Composing CNNs



# CNN example







# Typical CNN Layers

- **Convolutional Layer (CONV)**
- Pooling Layer (POOL)
- **Fully Connected (FC)**
- Normalization (NORM)



# Typical CNN Layers



- **Convolutional Layer (CONV)**
- Pooling Layer (POOL)
- **Fully Connected (FC)**
- Normalization (NORM)





# Typical CNN Layers

- Convolutional Layer (CONV)
- **Pooling Layer (POOL)**
- Fully Connected (FC)
- **Normalization (NORM)**  
Not commonly used







# Max Pooling and Average Pooling

4	9	2	5
5	6	2	4
2	4	5	4
5	6	8	4



9	5
6	8

**Max pool**

4	9	2	5
5	6	2	4
2	4	5	4
5	6	8	4



6.0	3.3
4.3	5.3

**Average pool**



# Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image


Max pool

Dim= 3 x 3  
Stride = 1

9		

Pooled Feature



# Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image


Max pool

Dim= 3 x 3  
Stride = 1

9	10	

Pooled Feature



# Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image


Max pool

Dim= 3 x 3  
Stride = 1

9	10	15

Pooled Feature



# Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image


Avg pool

Dim= 3 x 3  
Stride = 1

3		

Pooled Feature



# Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image


Avg pool

Dim= 3 x 3  
Stride = 1

3	4.8	

Pooled Feature



# Pooling Layers

1	2	5	10	7
0	5	1	4	0
0	4	9	3	15
0	0	2	4	5
0	7	2	0	0

Input Image


Avg pool


Dim= 3 x 3  
Stride = 1

3	4.8	6

Pooled Feature



# Full CNN example







# Benefits of CNNs

## **Parameter Sharing**

Filter can be useful in different parts of the input (image)

## **Sparsity of Connections**

- In each layer each output value depends only on small number of inputs (local)
- Translation invariance



# Challenges of CNNs

## **Computational Complexity**

Convolutions are expensive  $O(N^2n^4)$

## **Deeper Structure Needed**

In each layer each output value depends only on small number of inputs (local)



# Popular CNN Architectures (LeNet 5)



Yann LeCun




Leon Bottou



Yoshua Bengio




Patrick Haffner



Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 1998.




# LeNet-5 (1998)





# LeNet-5 (1998)



## Layer 1:

- Convolutional Layer with 6 kernels
- kernel size of 5x5
- Padding = 2, stride = 1

## Layer 2:

- Average pooling (2x2 kernel)

## Layer 3:

- Convolutional layer with 16 kernels
- kernel size of 5x5
- Padding = 0, stride = 1



# LeNet-5 (1998)

Layer 4:

Average pooling (2x2 kernel)

Layer 5:


- Convolutional layer with 120 kernels
- Kernel size of 5x5
- Padding = 0, stride = 1

Layer 6:

- Fully Connected Layer
- Input dimension = 120
- Output dimension = 84


Layer 7:

- Fully Connected Layer
- Input dimension = 84
- Output dimension = 10





# AlexNet (2012)





# AlexNet (2012)



Alex Krizhevsky



Ilya Sutskever




Geoffrey Hinton

Krizhevsky et al., Imagenet classification with deep convolutional neural networks, 2012






# Parameters (AlexNet)



Layer Name	Tensor Size	Weights	Biases	Parameters
Input Image	227x227x3	0	0	0
Conv-1	55x55x96	34,848	96	34,944
MaxPool-1	27x27x96	0	0	0
Conv-2	27x27x256	614,400	256	614,656
MaxPool-2	13x13x256	0	0	0
Conv-3	13x13x384	884,736	384	885,120
Conv-4	13x13x384	1,327,104	384	1,327,488
Conv-5	13x13x256	884,736	256	884,992
MaxPool-3	6x6x256	0	0	0
FC-1	4096x1	37,748,736	4,096	37,752,832
FC-2	4096x1	16,777,216	4,096	16,781,312
FC-3	1000x1	4,096,000	1,000	4,097,000
Output	1000x1	0	0	0
<b>Total</b>				<b>62,378,344</b>



# Parameters (AlexNet)



- Much bigger than LeNet (60M parameters)
- ReLU
- Multiple GPUs
- Local Response Normalization (LRN)



# VGG-16 (2014)

CONV:  $f=3$ ,  $s=1$ , same

POOL:  $f=2$ ,  $s=2$

Order: CCP CCP CCCP CCCP CCCP FFS

Nf:  $2^6$   $2^7$   $2^8$   $2^9$   $2^9$

~138 mil parameters





# VGG-16 (2014)

CONV:  $f=3$ ,  $s=1$ , same

POOL:  $f=2$ ,  $s=2$

Order: CCP CCP CCCP CCCP CCCP FFS

Nf:  $2^6$   $2^7$   $2^8$   $2^9$   $2^9$


~138 mil parameters

- Multiple convolution layers
- Smaller convolution filters
- Modularized architecture (VGG-19)





# Next episode in EEP 596...



## Recurrent Neural Networks