



LECTURE 4:

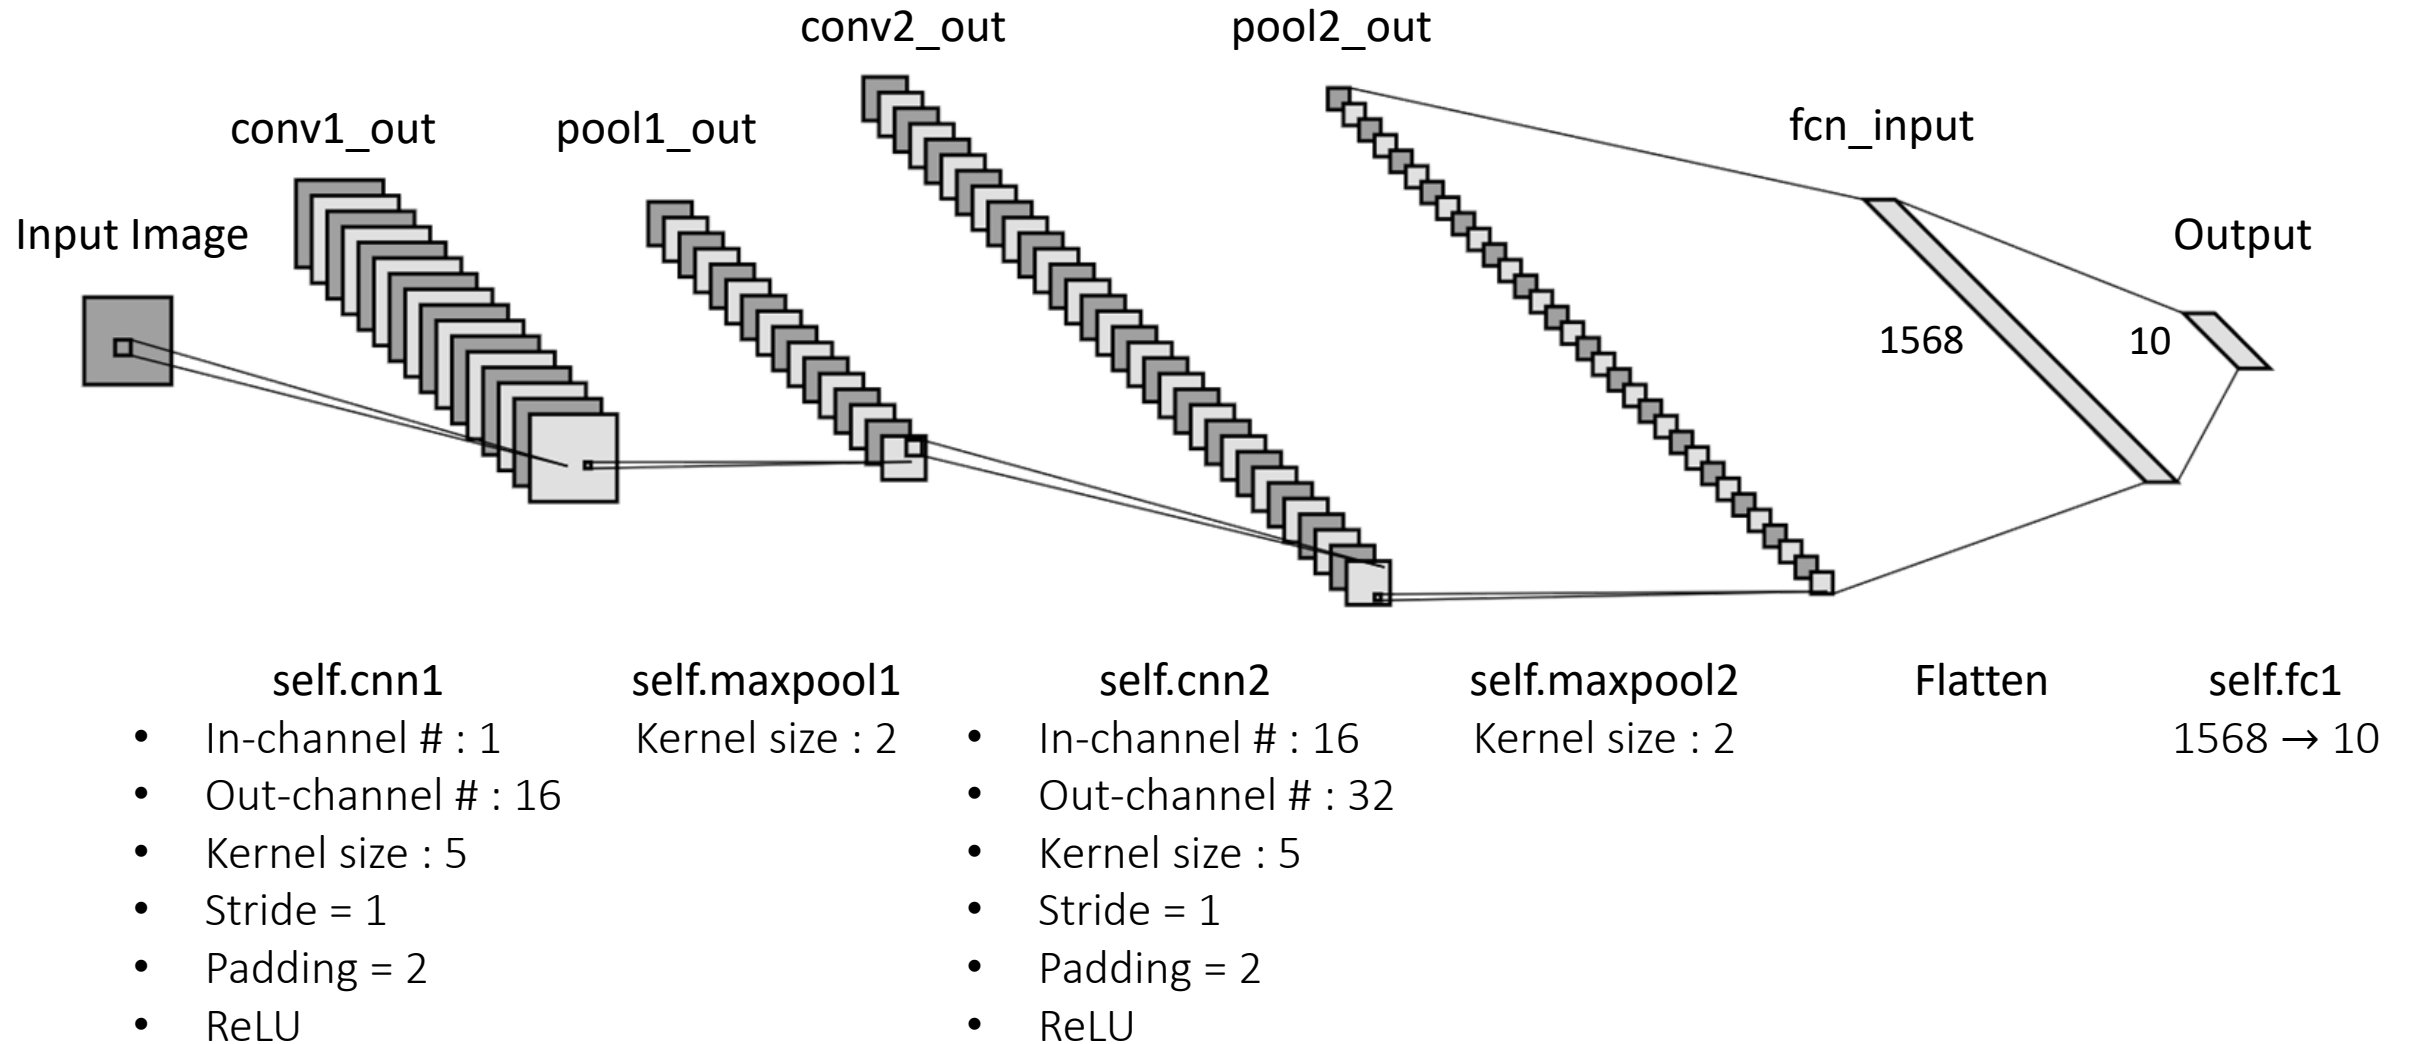
RECURRENT NEURAL NETWORK

University of Washington, Seattle

Fall 2024



Previously in EEP 596...





OUTLINE

Part 1: Introduction to RNNs

- Why do we need RNNs?
- RNN Architecture
- Embedding and Decoder

Part 3: RNN Problem Types

- RNN Configurations
- RNN Extensions

Part 2: Training RNNs

- Backpropagation in RNNs
- Vanishing/Exploding Gradient Problem
- Training with Teacher Forcing



INTRODUCTION TO RNNs


Why do we need RNNs?

RNN Architecture

Embedding and Decoder




Why Do We Need RNNs?



Feed-Forward Network




Why Do We Need RNNs?



Feed-Forward Network



Why Do We Need RNNs?



Feed-Forward Network has no memory of past inputs



Why Do We Need RNNs?

Korean

안녕하세요, 제 이름은 지민이에요



English

Hello, my name is Jimin



Why Do We Need RNNs?

Korean

안녕하세요,

English


Hello,



Why Do We Need RNNs?

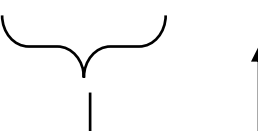
Korean

안녕하세요, 제



English

Hello, my

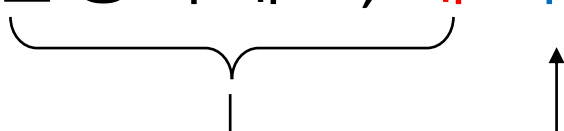




Why Do We Need RNNs?

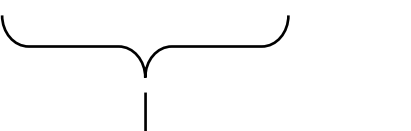
Korean

안녕하세요, 제 이름



English

Hello, my name

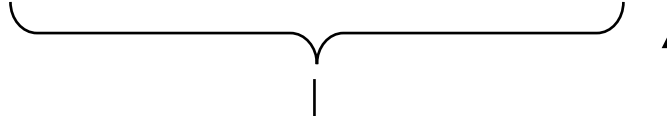




Why Do We Need RNNs?

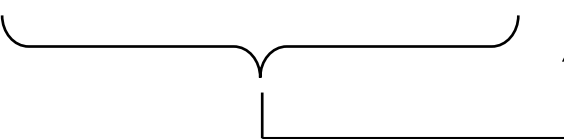
Korean

안녕하세요, 제 이름은



English

Hello, my name is





Why Do We Need RNNs?

Korean

안녕하세요, 제 이름은 지민이에요

The diagram illustrates the sequential nature of Korean. A bracket under '안녕하세요,' is followed by another bracket under '제 이름은', with an arrow pointing from the second bracket to '지민이에요', indicating that the meaning of the second part depends on the first.

English

Hello, my name is Jimin

The diagram illustrates the sequential nature of English. A bracket under 'Hello,' is followed by another bracket under 'my name is', with an arrow pointing from the second bracket to 'Jimin', indicating that the meaning of the second part depends on the first.



Why Do We Need RNNs?

Korean

안녕하세요, 제 이름은 지민이에요

English

Hello, my name is Jimin

Each word in a sentence is dependent to the past words →

Need memory



Why Do We Need RNNs?

Korean 안녕하세요, 제 이름은 지민이에요, 그리고 저는 비디오게임을 좋아해요

English Hello, my name is Jimin, and I like videogames

A sentence (input) could have **different sizes**



Why Do We Need RNNs?

We need a neural network architecture that can handle:



Why Do We Need RNNs?

We need a neural network architecture that can handle:

- Data order



Why Do We Need RNNs?

We need a neural network architecture that can handle:

- Data order
- Temporal dependencies



Why Do We Need RNNs?

We need a neural network architecture that can handle:


- Data order
- Temporal dependencies
- Variable input sizes



Why Do We Need RNNs?


We need a neural network architecture that can handle:

- Data order
- Temporal dependencies
- Variable input sizes





Brain is Highly Recurrent



Neurons themselves have continuous voltage dynamics

Different parts of brain exchange information both forward and backward



Brain is Highly Recurrent







RNN Architecture

Output

Hidden

Input



Unfold in Time





RNN Architecture

Output

Hidden

Input

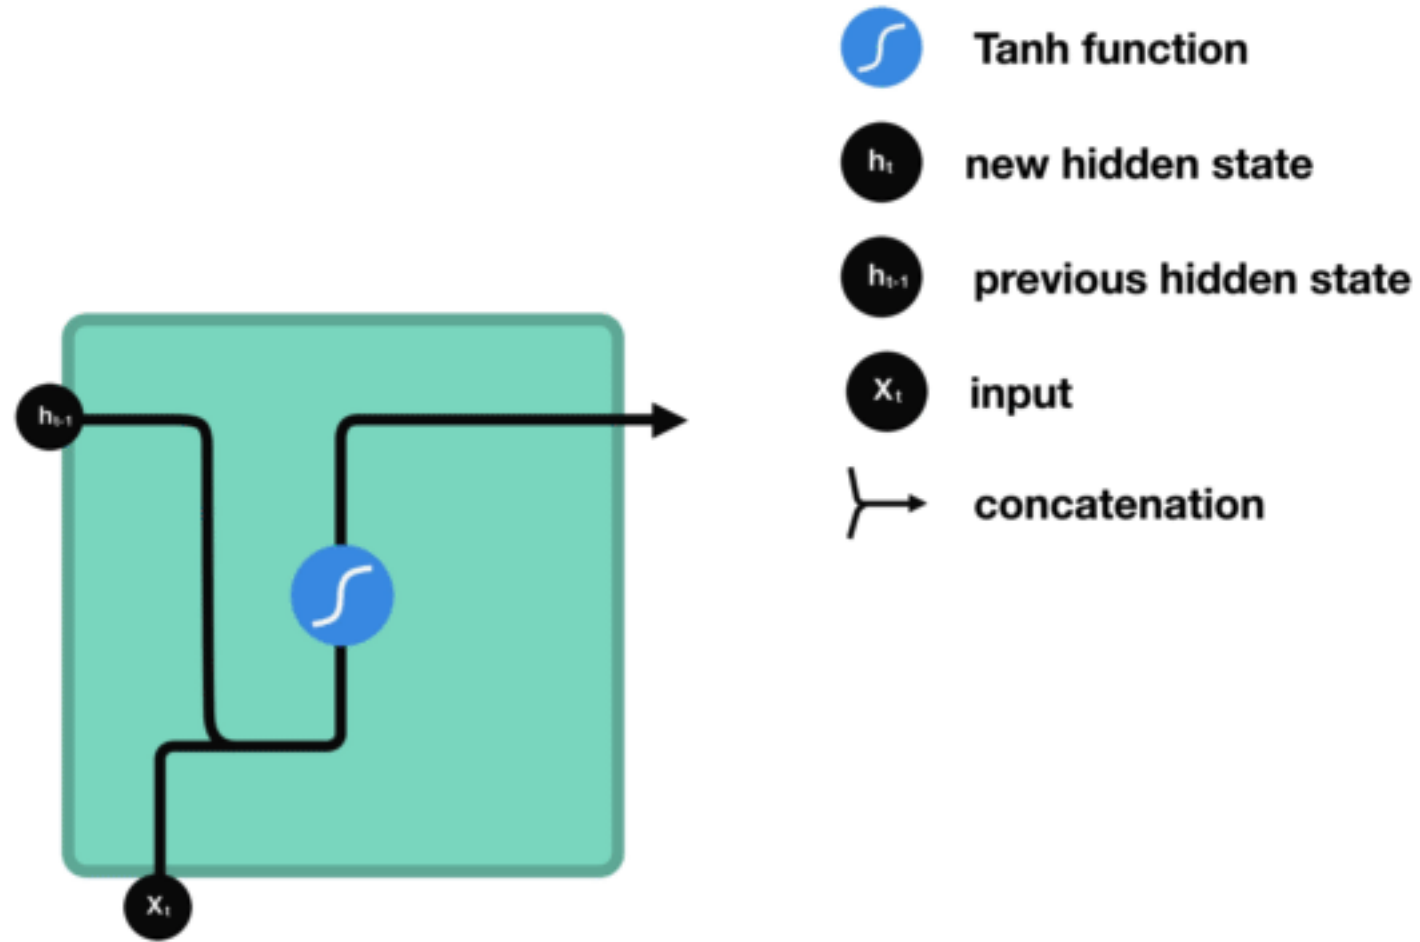


Unfold in Time

$$\begin{aligned} a^{(t)} &= b + \mathbf{W}h^{(t-1)} + \mathbf{U}x^{(t)} \\ h^{(t)} &= \tanh(a^{(t)}) \\ o^{(t)} &= c + \mathbf{V}h^{(t)} \\ \hat{y}^{(t)} &= \text{softmax}(o^{(t)}) \end{aligned}$$




RNN Architecture






RNN Architecture






RNN Architecture



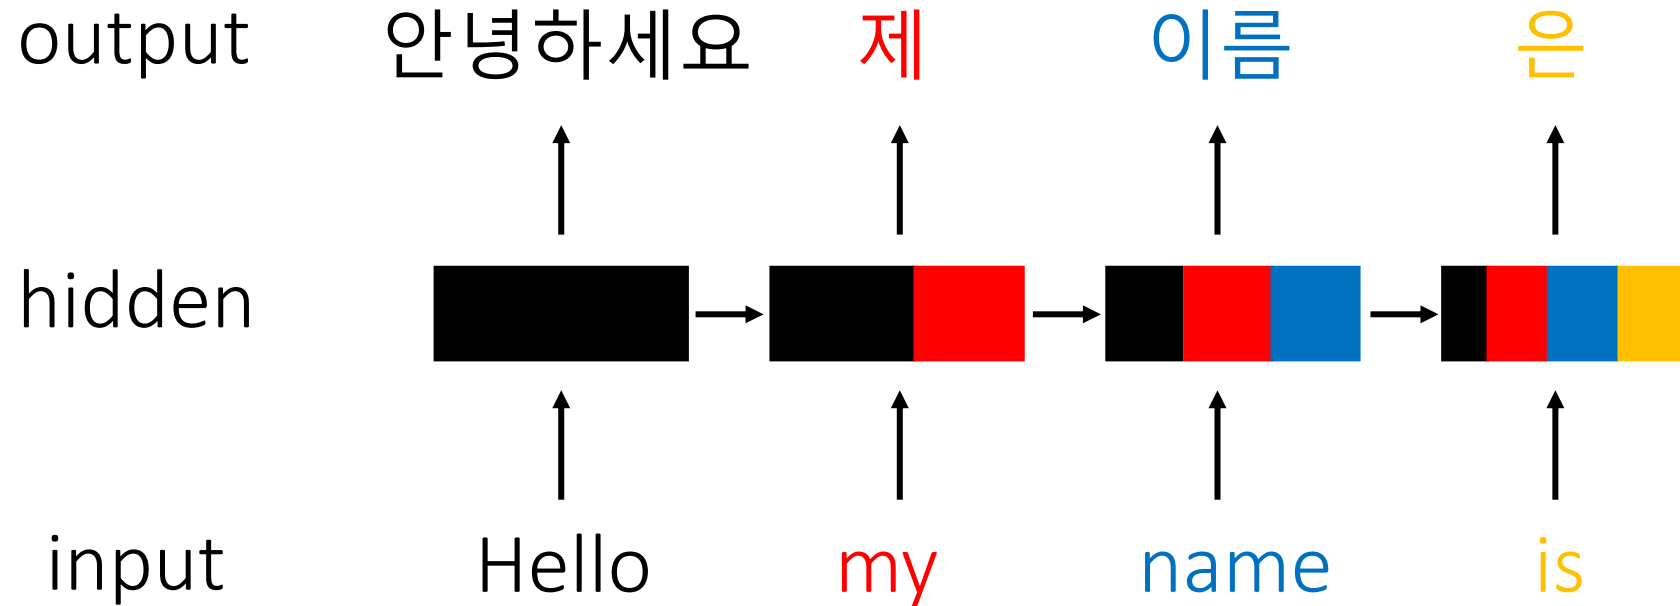


RNN Architecture






RNN Architecture





RNN Architecture





Sequential Data

Speech recognition



“The quick brown fox jumped
over the lazy dog.”

Music generation

∅



Sentiment classification

“There is nothing to like
in this movie.”



DNA sequence analysis

AGCCCCTGTGAGGAACTAG



AG**CCCCTGTGAGGAACT**AG

Machine translation

Voulez-vous chanter avec
moi?



Do you want to sing with
me?

Video activity recognition



Running

Name entity recognition


Yesterday, Harry Potter
met Hermione Granger.



Yesterday, **Harry Potter**
met **Hermione Granger**.

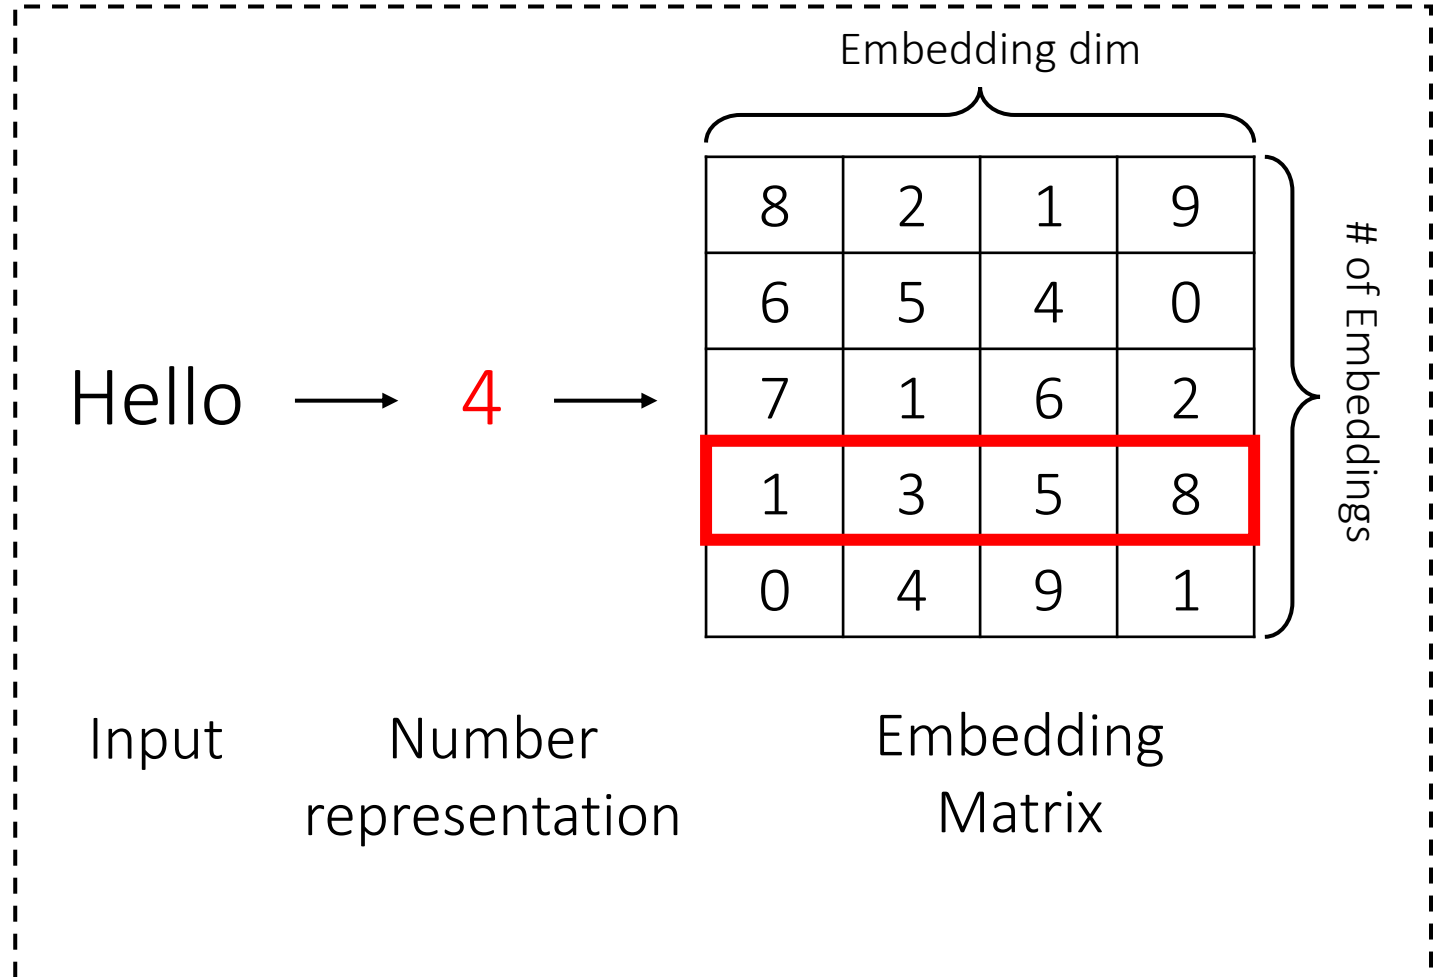



Embedding and Decoder



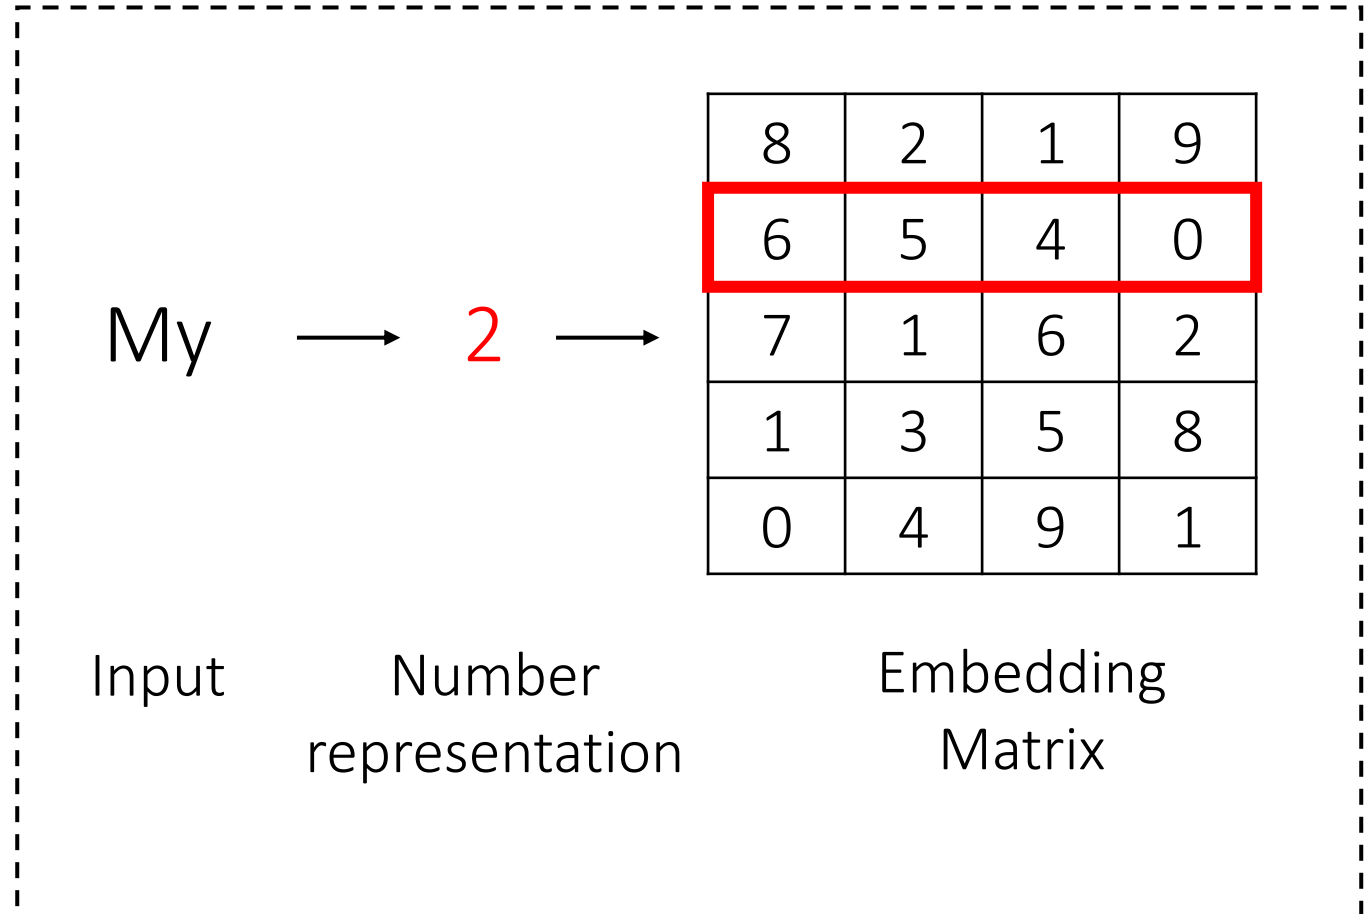



Embedding



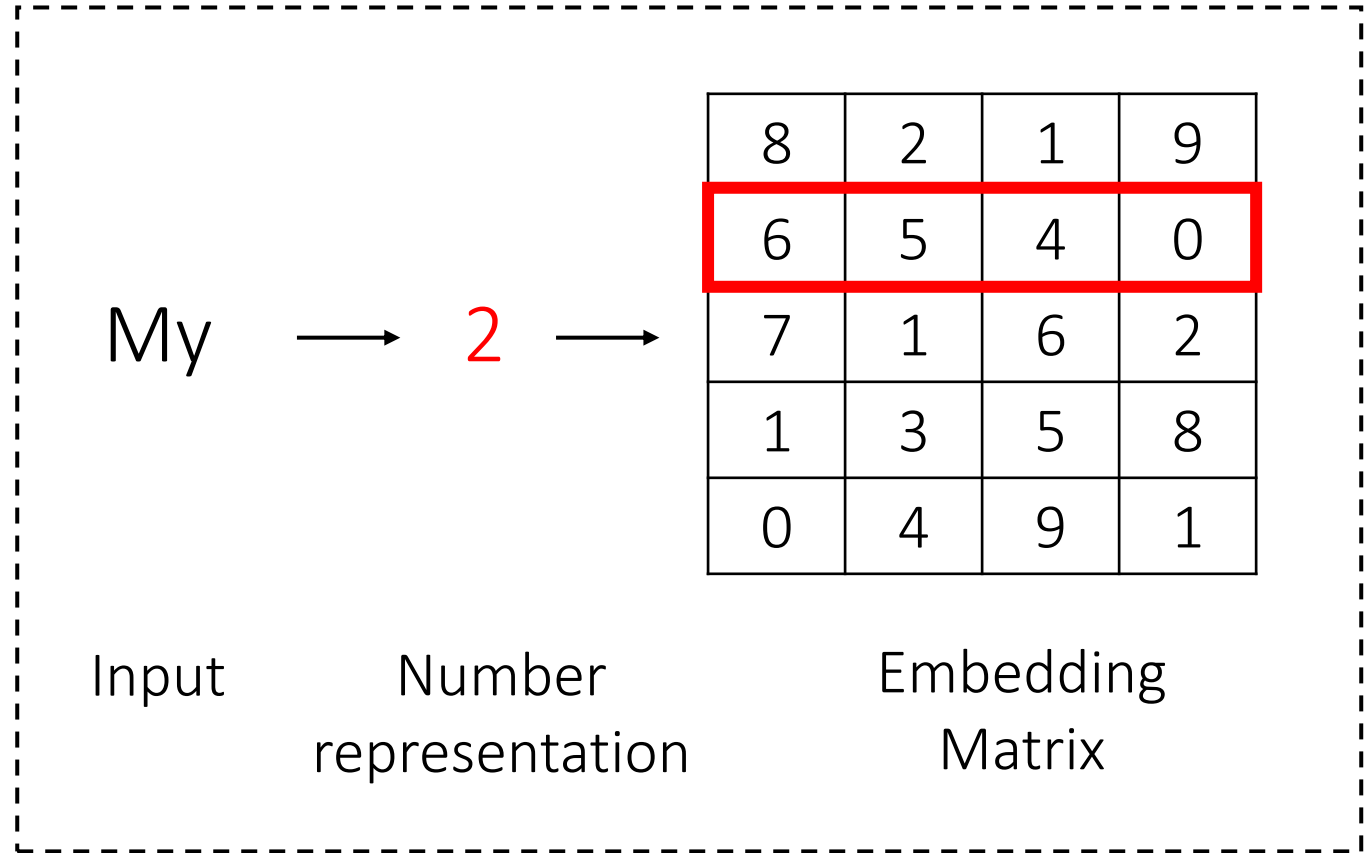



Embedding







Embedding



Embedding matrix is trainable



Embedding





`torch.nn.embedding(num_embeddings, embedding dim)`

<https://pytorch.org/docs/stable/generated/torch.nn.Embedding.html>



Decoder



```
torch.nn.Linear(hidden_size, output_size)
```



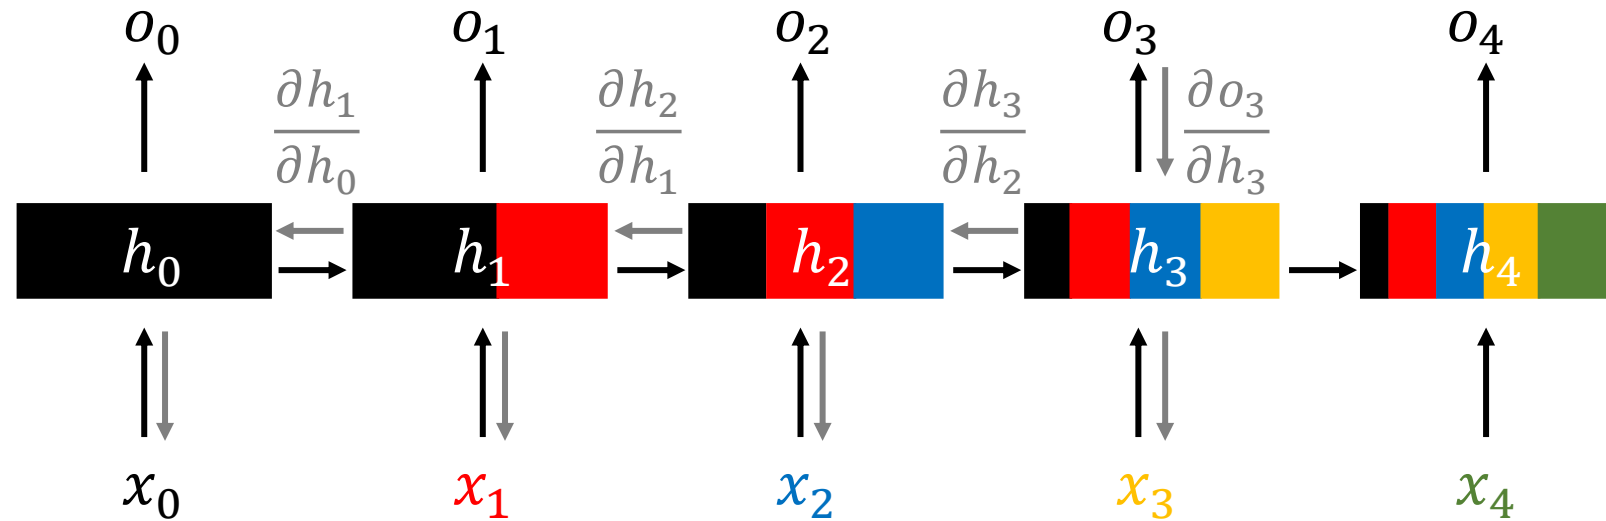
Backpropagation in RNNs

→ Forward
← Backward

output

hidden

input






Backpropagation in RNNs

→ Forward
← Backward

output

hidden

input



Backpropagation is performed backward in time


Vanishing and Exploding Gradients

→ Forward
← Backward

output

hidden

input




Vanishing and Exploding Gradients

→ Forward
← Backward

output

hidden

input




Longer input sequence →
higher risk of Vanishing/Exploding Gradients!



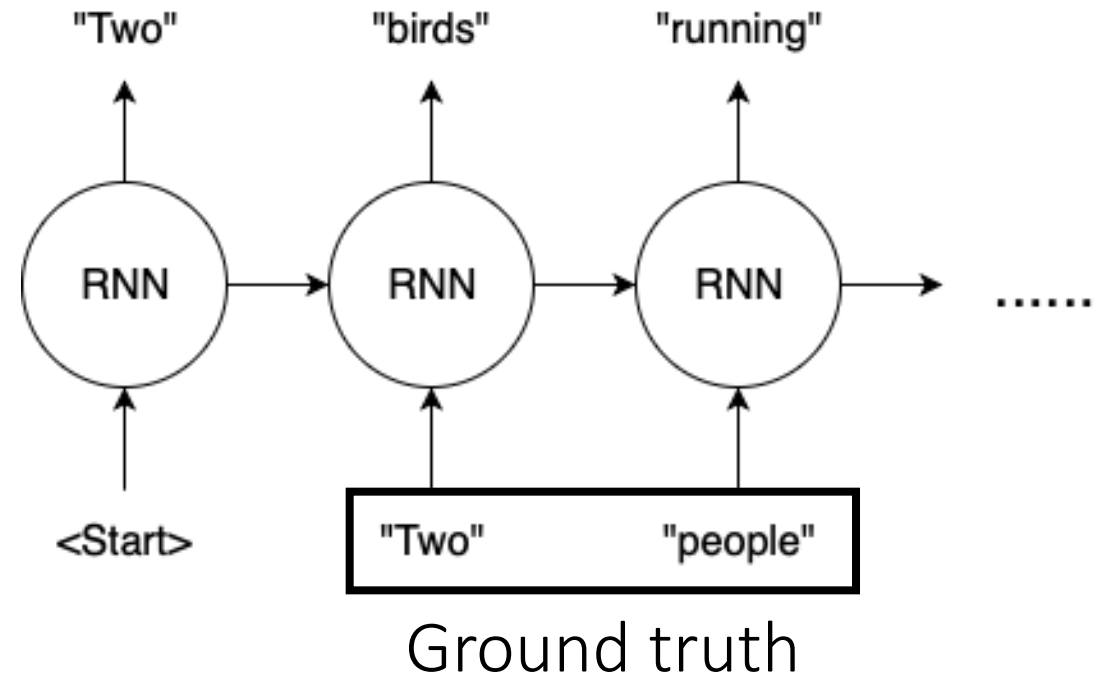
Vanishing and Exploding Gradients

- Use gated RNN architecture e.g., LSTM, GRU (Lab 6)
- ReLU activation as nonlinearity
- Smaller number of sequence
- Smaller learning rate

Training RNN with Teacher Forcing



Without Teacher Forcing



With Teacher Forcing



RNN PROBLEM TYPES


RNN Configurations

RNN Extensions




RNN Configurations


one to one




one to many




many to one



many to many




many to many





One to One


one to one





One to One

one to one




Identical to Feed Forward Network



One to Many

one to many





One to Many


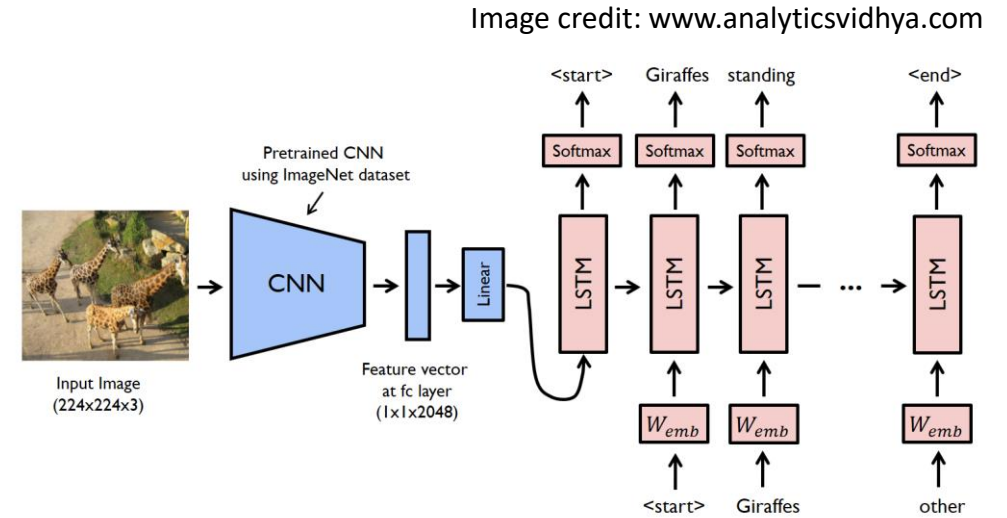



Image captioning




Music generation



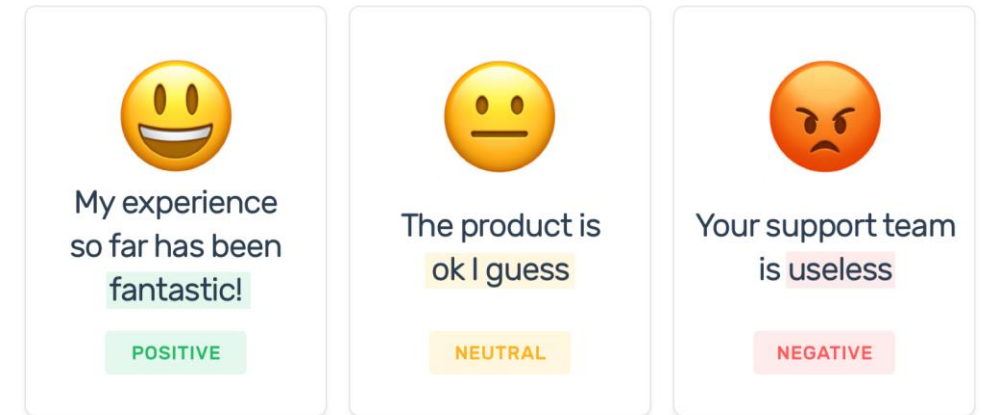
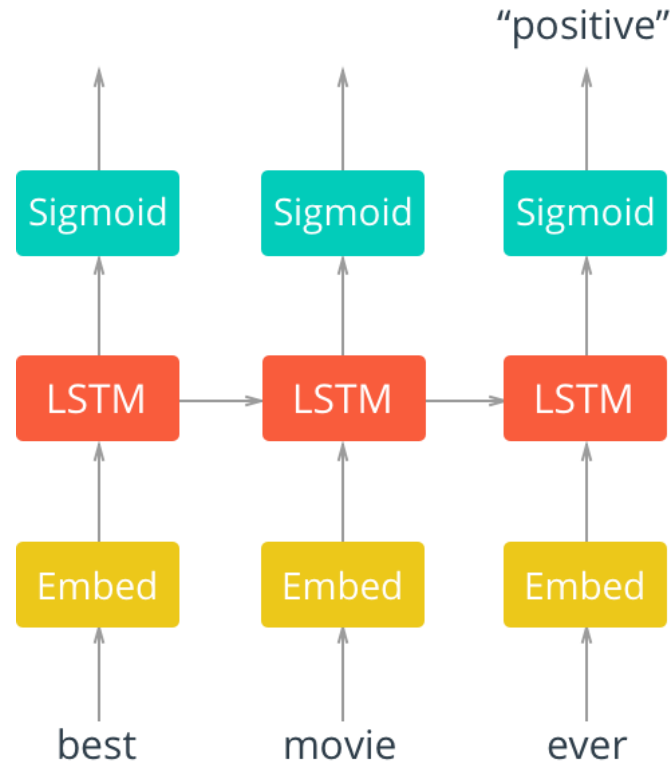



Many to One





Many to One




Sentiment Analysis




Many to Many

many to many




many to many







Many to Many

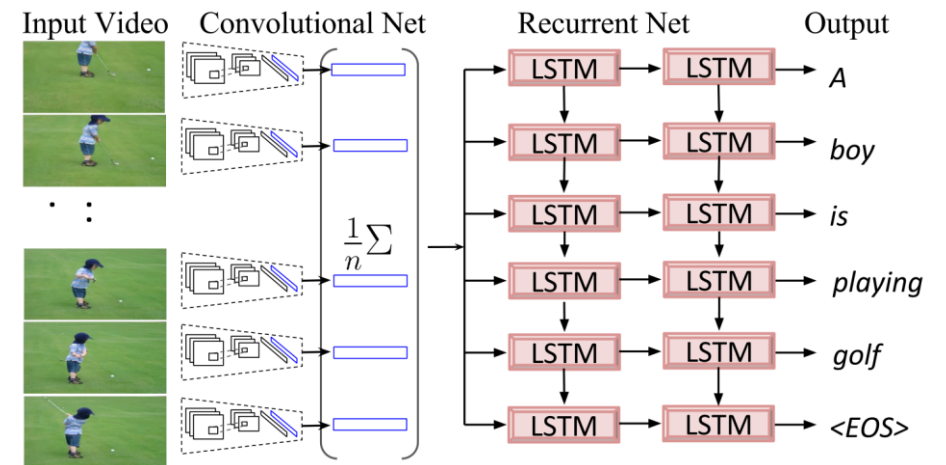
many to many



many to many




Machine Translation




Video Captioning




RNN Extensions



Regular RNN




Deep RNN



Bi-directional RNN



Deep RNN



Deep RNN

(+)

Can provide better performance
Often used for complex problems


(-)

Potential for overfitting
Longer training time

Implemented as 'num_layers'
parameters in `torch.nn.RNN()`



Bi-directional RNNs



Bi-directional RNN

(+)

Higher performance in Natural Language Processing tasks

Suitable when both left and right contexts are used

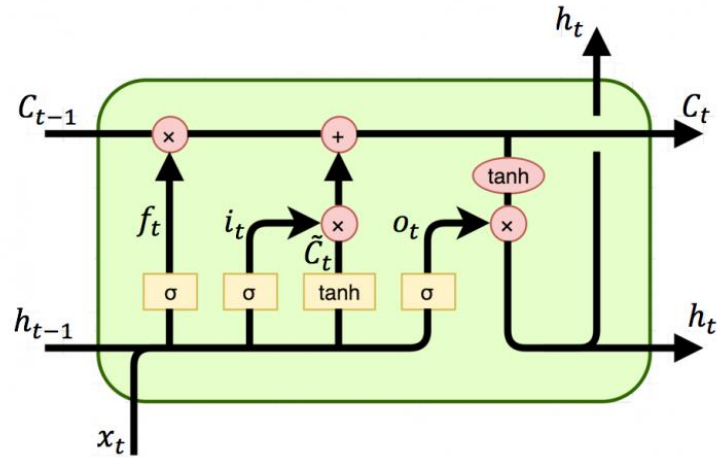
(-)

Harder to train than Uni-directional RNN
Not suitable for real-time processing


Implemented as 'bidirectional' parameter in `torch.nn.RNN()`



Next episode in EEP 596...



Long-short-term memory
(LSTM)



Gated recurrent units
(GRU)



Next episode in EEP 596...

