



# LAB 6: ADVANCED RECURRENT NEURAL NETWORKS

University of Washington, Seattle

Fall 2024



# OUTLINE

## Part 1: Gated RNNs

- Need for Gated RNNs
- LSTM
- GRU

## Part 2: Training Gated RNNs

- Mini-batch Gradient in RNNs
- RNN extensions on LSTM/GRU

## Part 3: Gated RNN Application with PyTorch

- Signal Denoising

## Part 4: Encoder-Decoder RNNs in PyTorch

- Signal Prediction

## Part 5: Lab Assignment

- Stock Prediction



# GATED RNNs


Need for Gated RNNs

Long Short-Term Memory (LSTM)

Gated Recurrent Unit (GRU)



# Recap: Backpropagation in RNNs






# Recap: Backpropagation in RNNs

→ Forward  
← Backward

output

hidden

input



Backpropagation is performed backward in time


# Vanishing and Exploding Gradients

→ Forward  
← Backward

output

hidden

input




# Vanishing and Exploding Gradients

→ Forward  
← Backward

output

hidden

input



Longer input sequence →  
higher risk of Vanishing/Exploding Gradients!


# Vanishing and Exploding Gradients

→ Forward  
← Backward

output

hidden

input




Need for better RNN architecture capable of  
processing longer sequence

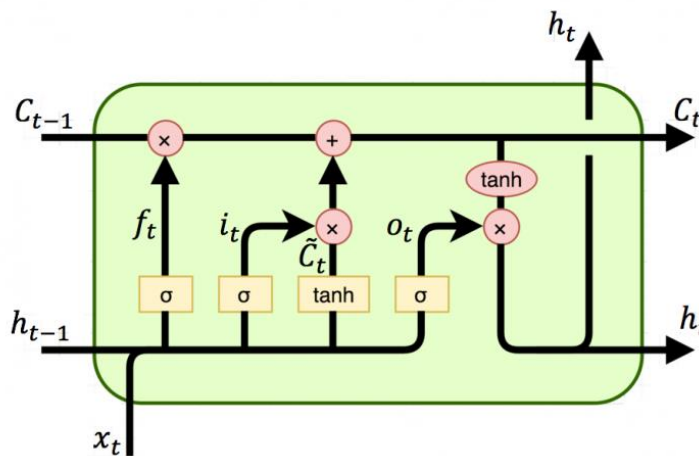





# Gated RNNs




Vanilla RNN



LSTM




GRU





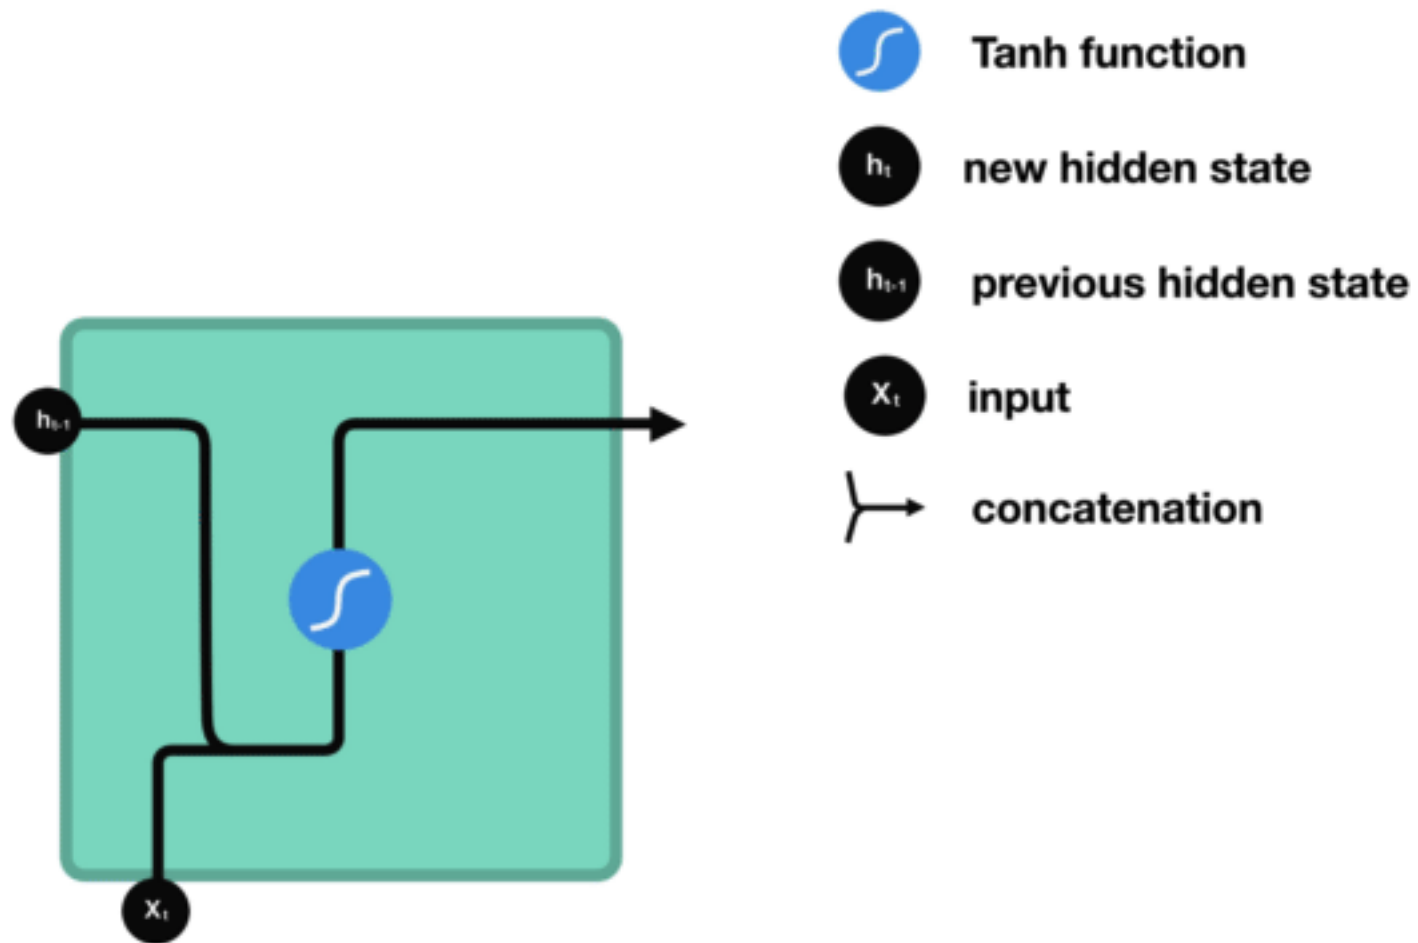
# Vanilla RNN



Vanilla RNN




# Vanilla RNN





# Vanilla RNN



Tanh function



new hidden state



previous hidden state



input



concatenation

$$\mathbf{a}^{(t)} = \mathbf{b} + \mathbf{W}\mathbf{h}^{(t-1)} + \mathbf{U}\mathbf{x}^{(t)}$$

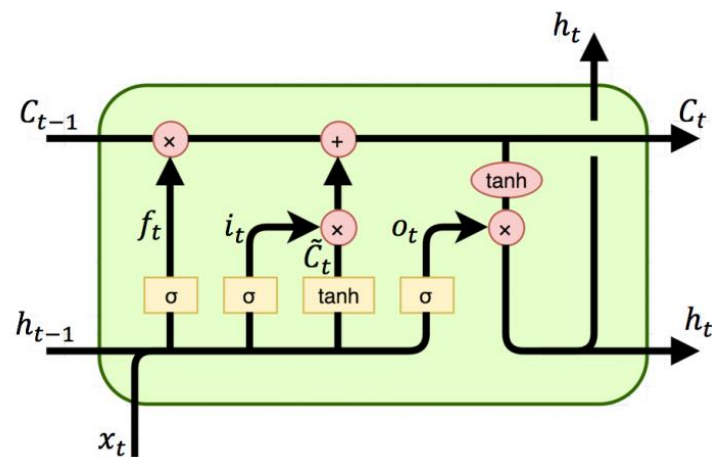
$$\mathbf{h}^{(t)} = \tanh(\mathbf{a}^{(t)})$$

$$\mathbf{o}^{(t)} = \mathbf{c} + \mathbf{V}\mathbf{h}^{(t)}$$


$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{o}^{(t)})$$



# LSTM (Long Short-Term Memory)

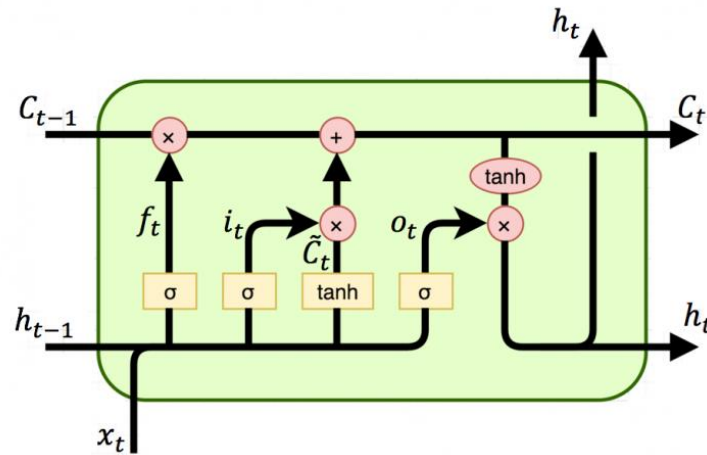


LSTM





# LSTM (Long Short-Term Memory)



LSTM

$$f_t = \sigma_g(W_f x_t + U_f h_{t-1} + b_f)$$

$$i_t = \sigma_g(W_i x_t + U_i h_{t-1} + b_i)$$

$$o_t = \sigma_g(W_o x_t + U_o h_{t-1} + b_o)$$

$$c_t = f_t \circ c_{t-1} + i_t \circ \sigma_c(W_c x_t + U_c h_{t-1} + b_c)$$

$$h_t = o_t \circ \sigma_h(c_t)$$



Neural Network  
Layer



Pointwise  
Operation



Vector  
Transfer




Concatenate



Copy



# LSTM: Detailed Architecture




## Cell state

- Unique to LSTM
- Long term memory of the model

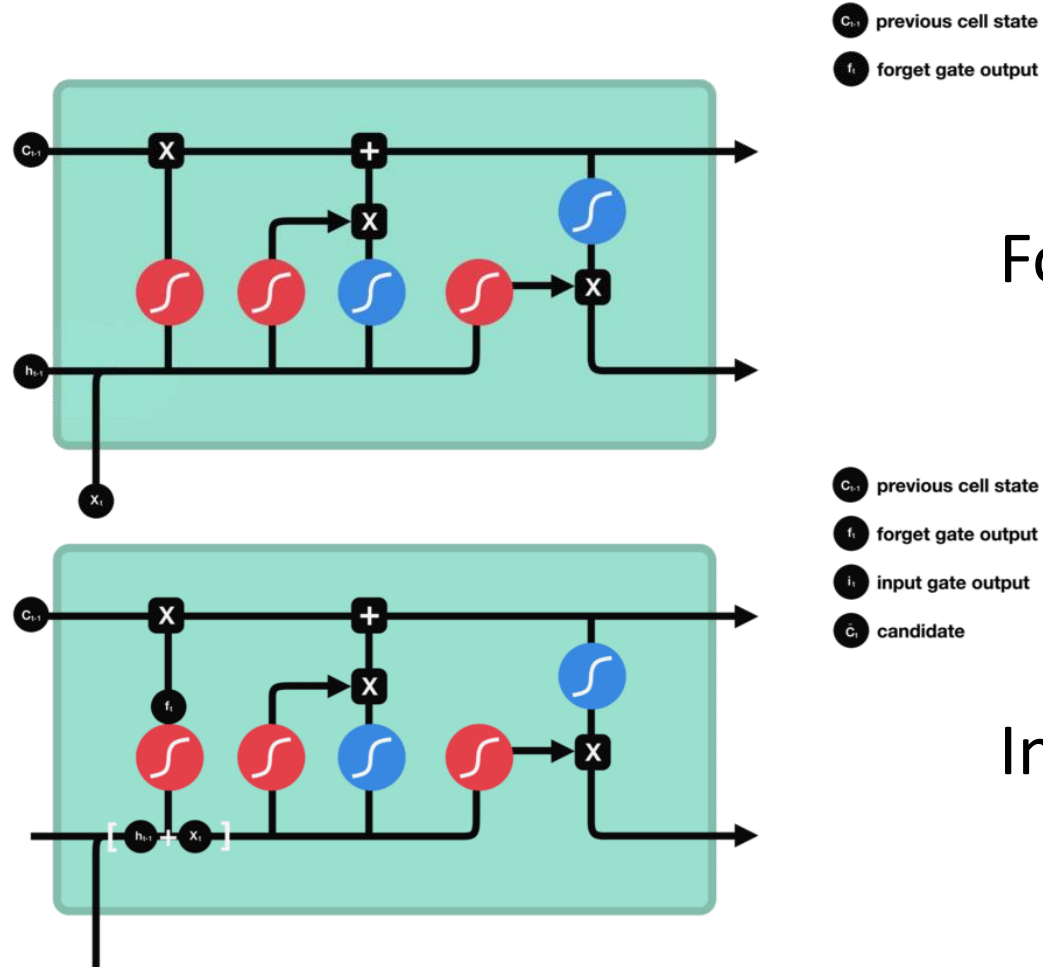



# LSTM: Detailed Architecture

$$f_t = \sigma_g(W_f x_t + U_f h_{t-1} + b_f)$$



$$i_t = \sigma_g(W_i x_t + U_i h_{t-1} + b_i)$$



Forget gate layer


Input gate layer







# LSTM: Detailed Architecture

$$c_t = f_t \circ c_{t-1} + i_t \circ \sigma_c(W_c x_t + U_c h_{t-1} + b_c)$$




$$o_t = \sigma_g(W_o x_t + U_o h_{t-1} + b_o)$$



$$c_t = f_t \cdot c_{t-1} + i_t \cdot \tilde{c}_t$$

- $c_{t-1}$  previous cell state
- $f_t$  forget gate output
- $i_t$  input gate output
- $\tilde{c}_t$  candidate
- $c_t$  new cell state

Update cell state



- $c_{t-1}$  previous cell state
- $f_t$  forget gate output
- $i_t$  input gate output
- $\tilde{c}_t$  candidate
- $c_t$  new cell state
- $o_t$  output gate output
- $h_t$  hidden state

Output gate layer



# LSTM: Detailed Architecture

## Forget gate

Decides what is relevant to keep from previous steps

## Input gate


Decides what information is relevant to add from the current step

## Output Gate


Determines what the next hidden state should be



# Gated RNNs




GRU





# GRU: Detailed Architecture



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$


$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$



# Information Flow in GRU





# GRU: Detailed Architecture

## Update gate

How much of the past information needs to be retained

## Reset gate

How much of the past information to forget




# TRAINING GATED RNNs

Mini-batch Gradient in RNNs

RNN Extensions in LSTM/GRU



# Mini-batch Gradient in RNNs



RNN input format in PyTorch = (batch size, sequence length, # of features)


Example above = (4, 17, 1)

with batch\_first = True







# Gated RNNs



Vanilla RNN



LSTM




GRU

Inputs =  $x_t$


Outputs =  $f(h(t))$




# RNN Extensions in LSTM/GRU



Regular RNN



Deep RNN



Bi-directional RNN



# RNN Extensions in LSTM/GRU

```
1 class example_LSTM(torch.nn.Module):
2
3     def __init__(self, input_size, hidden_size, num_layers, output_size):
4
5         super(example_LSTM, self).__init__()
6
7         self.lstm = torch.nn.LSTM(input_size=input_size, hidden_size=hidden_size,
8                                   num_layers = num_layers,
9                                   batch_first = True,
10                                  bidirectional = False,
11                                  dropout = 0.1)
12
13         self.decoder = torch.nn.Linear(hidden_size, output_size)
14
15     def forward(self, input_seq, hidden_state):
16
17         pred, hidden = self.lstm(input_seq, hidden_state)
18
19         pred = self.decoder(pred)
20
21         return pred
```

**num\_layers:**  
LSTM layers to be stacked

**batch\_first:**  
Tells PyTorch we are using  
(batchsize, seq\_len, feature #)

**bidirectional:**  
Whether to configure  
bidirectional LSTM

**dropout:**  
introduces dropout layer on the  
outputs of each LSTM layer  
except for last layer  
(use when num\_layers > 1)



# RNN Extensions in LSTM/GRU

```
1 class example_GRU(torch.nn.Module):
2
3     def __init__(self, input_size, hidden_size, num_layers, output_size):
4
5         super(example_GRU, self).__init__()
6
7         self.gru = torch.nn.GRU(input_size=input_size, hidden_size=hidden_size,
8                                 num_layers = num_layers,
9                                 batch_first = True,
10                                bidirectional = False,
11                                dropout = 0.1)
12
13         self.decoder = torch.nn.Linear(hidden_size, output_size)
14
15     def forward(self, input_seq, hidden_state):
16
17         pred, hidden = self.gru(input_seq, hidden_state)
18
19         pred = self.decoder(pred)
20
21         return pred
```

**num\_layers:**  
GRU layers to be stacked

**batch\_first:**  
Tells PyTorch we are using  
(batchsize, seq\_len, feature #)

**bidirectional:**  
Whether to configure  
bidirectional GRU

**dropout:**  
introduces dropout layer on the  
outputs of each GRU layer except  
for last layer  
(use when num\_layers > 1)





# IMPLEMENTATION OF GATED RNNs in PYTORCH

Signal Denoising




# Signal Denoising

Output Sequence



Input Sequence





# Prepare Data

```
1 def sinusoidal_generator(X, signal_freq=60.):
2     return np.sin(2 * np.pi * (X) / signal_freq)
3
4
5 def add_noise(Y, noise_range=(-0.35, 0.35)):
6
7     noise = np.random.uniform(noise_range[0], noise_range[1], size=Y.shape)
8
9     return Y + noise
10
11 def sample_seq(sequence_length):
12
13     random_offset = random.randint(0, sequence_length)
14     X = np.arange(sequence_length)
15
16     denoised_output_seq = sinusoidal_generator(X + random_offset)
17     noisy_input_seq = add_noise(denoised_output_seq)
18
19     return noisy_input_seq, denoised_output_seq
```

Sinusoidal wave generator

Add noise function

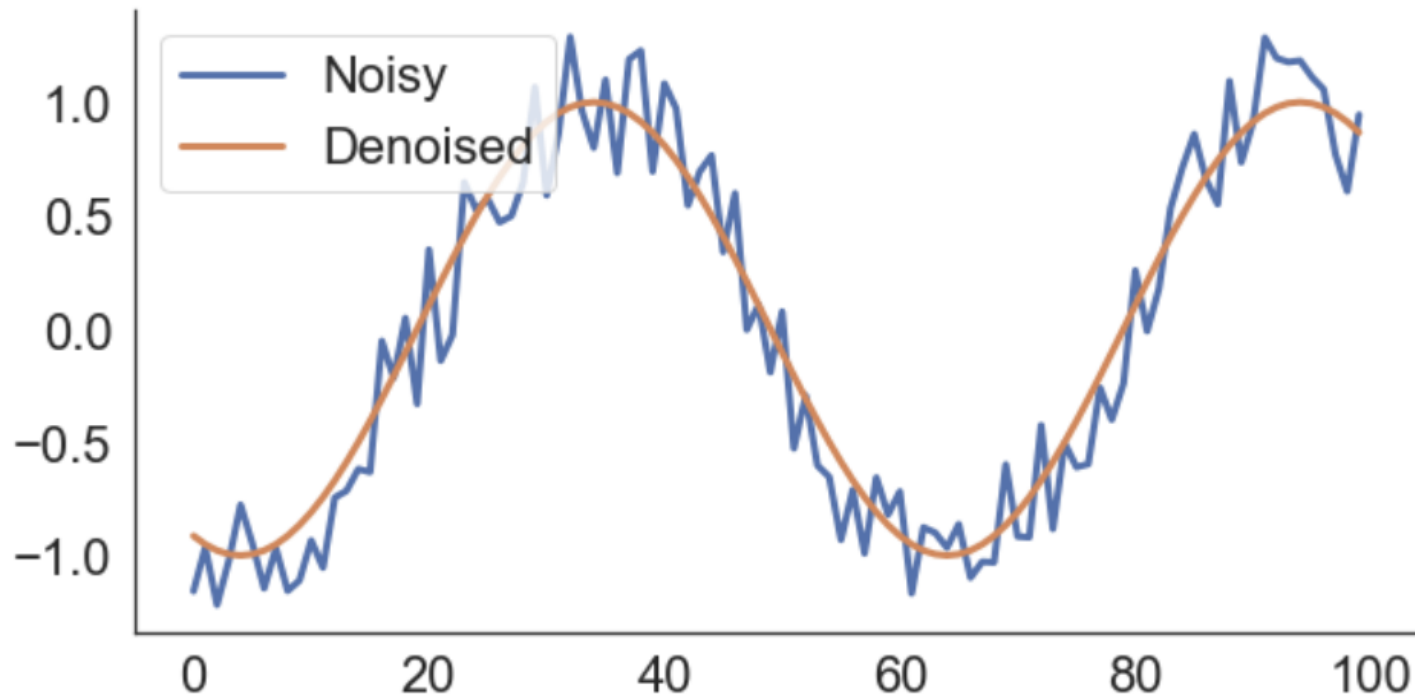
Generate sample ground truth/noisy  
sinusoidal waves



# Prepare Data

```
1 noisy_input_seq, denoised_output_seq = sample_seq(sequence_length = 100)
2
3 plt.figure(figsize = (10, 5))
4
5 plt.plot(noisy_input_seq, label = 'Noisy', linewidth = 3)
6 plt.plot(denoised_output_seq, label = 'Denoised', linewidth = 3)
7 plt.legend()
8 sns.despine()
```

Example sample ground truth & noisy sinusoidal wave with sequence length = 100







# Prepare Data

```
1 def create_synthetic_dataset(n_samples, sequence_length):
2
3     noisy_seq_inputs = np.zeros((n_samples, sequence_length))
4     denoised_seq_outputs = np.zeros((n_samples, sequence_length))
5
6     for i in range(n_samples):
7
8         noisy_inp, denoised_out = sample_seq(sequence_length)
9
10        noisy_seq_inputs[i, :] = noisy_inp
11        denoised_seq_outputs[i, :] = denoised_out
12
13    return noisy_seq_inputs, denoised_seq_outputs
```

Using the sample\_seq() function to generate synthetic ground truth/noisy dataset of n-samples

```
1 noisy_seq_inputs, denoised_seq_outputs = create_synthetic_dataset(n_samples = 12000,
2                                                                    sequence_length = 100)
3
4 train_input_seqs, train_output_seqs = noisy_seq_inputs[:8000], denoised_seq_outputs[:8000]
5 test_input_seqs, test_output_seqs = noisy_seq_inputs[8000:], denoised_seq_outputs[8000:]
```

Take first 8000 as training dataset and 4000 as testing dataset

```
1 train_input_seqs = train_input_seqs.reshape((train_input_seqs.shape[0], -1, 1))
2 train_output_seqs = train_output_seqs.reshape((train_output_seqs.shape[0], -1, 1))
3
4 test_input_seqs = test_input_seqs.reshape((test_input_seqs.shape[0], -1, 1))
5 test_output_seqs = test_output_seqs.reshape((test_output_seqs.shape[0], -1, 1))
```

Reshape training and testing dataset to conform to (# of samples, seq\_len, feature #) format



# Define Model

```
1 class Denoiser_GRU(torch.nn.Module):
2
3     def __init__(self, input_size, hidden_size, num_layers, output_size):
4
5         super(Denoiser_GRU, self).__init__()
6
7         self.gru = torch.nn.GRU(input_size=input_size, hidden_size=hidden_size,
8                                 num_layers = num_layers,
9                                 batch_first = True,
10                                bidirectional = False)
11
12         self.decoder = torch.nn.Linear(hidden_size, output_size)
13
14         self.output_activation = torch.nn.Tanh()
15
16     def forward(self, input_seq, hidden_state):
17
18         pred, hidden = self.gru(input_seq, hidden_state)
19
20         pred = self.output_activation(self.decoder(pred))
21
22         return pred
```

Using GRU with batch\_first = True

Decoder layer to convert hidden states to final output

Using **Tanh** on decoder output layer to squeeze output value between -1 and 1

Input\_sequence, hidden\_states → GRU →  
output\_sequence, hidden\_states →  
Decoder Layer → Tanh activation



# Define Hyperparameters

```
1 denoiser_GRU = Denoiser_GRU(input_size = 1, hidden_size = 30,  
2                             num_layers = 1, output_size = 1)  
3  
4 learning_rate = 0.0003  
5 epochs = 100  
6  
7 batchsize = 300  
8  
9 loss_func = torch.nn.L1Loss()  
10 optimizer = torch.optim.Adam(denoiser_GRU.parameters(), lr=learning_rate)  
11  
12 denoiser_GRU
```

Input dim to GRU = 1

Hidden state size = 30

GRU layers to be stacked = 1

Output dim of decoder layer = 1

Define learning rate, epochs and batch size

Using L1Loss (Least Absolute Deviations) and Adam optimizer



# Identify Tracked Values

```
1 train_loss_list = []
```

Empty Python list to keep track of training loss



# Train Model

```
1 train_input_seqs = torch.from_numpy(train_input_seqs).float()
2 train_output_seqs = torch.from_numpy(train_output_seqs).float()
3
4 test_input_seqs = torch.from_numpy(test_input_seqs).float()
5 test_output_seqs = torch.from_numpy(test_output_seqs).float()
6
7 train_batches_features = torch.split(train_input_seqs, batchsize)
8 train_batches_targets = torch.split(train_output_seqs, batchsize)
9
10 batch_split_num = len(train_batches_features)
11
12 for epoch in range(epochs):
13     for k in range(batch_split_num):
14         hidden_state = None
15
16         pred = denoiser_GRU(train_batches_features[k], hidden_state)
17
18         optimizer.zero_grad()
19
20         loss = loss_func(pred, train_batches_targets[k])
21         train_loss_list.append(loss.item())
22
23         loss.backward()
24
25         optimizer.step()
26
27
28
29 print("Averaged Training Loss for Epoch ", epoch, ": ", np.mean(train_loss_list[-batch_split_num:]))
```

Convert training and testing data to Tensors

Split training data into mini-batches


Training loop with mini-batch gradient

Print averaged loss throughout the epoch



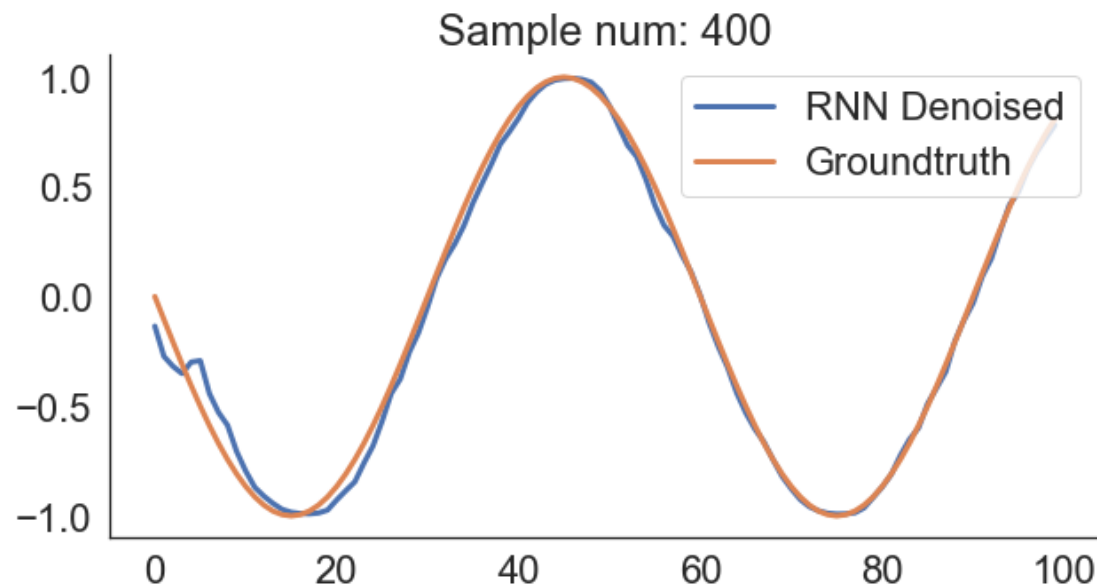
# Visualize & Evaluate Model

```
1 plt.figure(figsize = (10, 5))
2
3 plt.plot(train_loss_list, linewidth = 3, label = 'Training Loss')
4 plt.ylabel("training loss")
5 plt.xlabel("Iterations")
6 plt.legend()
7 sns.despine()
```



```
1 with torch.no_grad():
2
3     test_prediction = denoiser_GRU(test_input_seqs, None)
4     print("Testing Loss (Least Absolute Deviations): ",
5           loss_func(test_prediction, test_output_seqs).item())
```

Testing Loss (Least Absolute Deviations): 0.04158925637602806






# ENCODER-DECODER APPLICATION IN PYTORCH

Signal Prediction





# Example Task Description







# Example Task Description



Encoder input sequence length = 5

Decoder output sequence length = 2  
(prediction)



# Prepare Data

```
1 def generate_noisy_signal(datapoints_num, tf):
2
3     t = np.linspace(0., tf, datapoints_num)
4     y = np.sin(2. * t) + 0.5 * np.cos(t) + np.random.normal(0., 0.2, datapoints_num)
5
6     return y.reshape(-1, 1)
```

Function for generating a noisy signal  
(sin + cos + noise)

```
1 def generate_input_output_seqs(y, encoder_inputseq_len, decoder_outputseq_len, stride = 1, num_features = 1):
2
3     L = y.shape[0]
4     num_samples = (L - encoder_inputseq_len - decoder_outputseq_len) // stride + 1
5
6     train_input_seqs = np.zeros([num_samples, encoder_inputseq_len, num_features])
7     train_output_seqs = np.zeros([num_samples, decoder_outputseq_len, num_features])
8
9     for ff in np.arange(num_features):
10
11         for ii in np.arange(num_samples):
12
13             start_x = stride * ii
14             end_x = start_x + encoder_inputseq_len
15             train_input_seqs[ii, :, ff] = y[start_x:end_x, ff]
16
17             start_y = stride * ii + encoder_inputseq_len
18             end_y = start_y + decoder_outputseq_len
19             train_output_seqs[ii, :, ff] = y[start_y:end_y, ff]
20
21     return train_input_seqs, train_output_seqs
```

Function for generating

- input sequences to encoder
- output target sequences for decoder

e.g.,  $y = [1, 2, 3, 4, 5, 6, 7, 8]$

Encoder inputseq len = 3

Decoder outputseq len = 2

train\_input\_seqs =  
[[1,2,3],[2,3,4],[3,4,5],[4,5,6]]  
train\_output\_seqs =  
[[4,5],[5,6],[6,7],[7,8]]



# Prepare Data

```
1 encoder_inputseq_len = 5
2 decoder_outputseq_len = 2
3 testing_sequence_len = 50
4
5 y = generate_noisy_signal(datapoints_num = 2000, tf = 80 * np.pi)
6 y_train = y[:-testing_sequence_len]
```

- Encoder input sequence length = 5
- Decoder output sequence length = 2
- Testing sequence length = 50


```
1 train_input_seqs, train_output_seqs = generate_input_output_seqs(y = y_train,
2                                                                    encoder_inputseq_len = encoder_inputseq_len,
3                                                                    decoder_outputseq_len = decoder_outputseq_len,
4                                                                    stride = 1,
5                                                                    num_features = 1)
```

```
1 print("Encoder Training Inputs Shape: ", train_input_seqs.shape)
2 print("Decoder Training Outputs Shape: ", train_output_seqs.shape)
```

Encoder Training Inputs Shape: (1944, 5, 1)  
Decoder Training Outputs Shape: (1944, 2, 1) (sample size, sequence length, feature/timestep)



# Prepare Data




`train_input_seqs[0]`  
(input to encoder)

`train_output_seqs[0]`  
(output target by decoder)



# Prepare Data




train\_input\_seqs[1]

train\_output\_seqs[1]




# Prepare Data






# Prepare Data





# Define Model



Using LSTM for Encoder


No need for FC layer since encoder only passes hidden states to Decoder

```
1 class Encoder(torch.nn.Module):
2
3     def __init__(self, input_size, hidden_size, num_layers):
4
5         super(Encoder, self).__init__()
6
7         self.lstm = torch.nn.LSTM(input_size = input_size, hidden_size = hidden_size,
8                                   num_layers = num_layers,
9                                   batch_first = True)
10
11     def forward(self, input_seq, hidden_state):
12
13         lstm_out, hidden = self.lstm(input_seq, hidden_state)
14
15         return lstm_out, hidden
```





# Define Model




Using LSTM for Decoder

FC layer for converting hidden states to a single number (prediction)

```
17 class Decoder(torch.nn.Module):
18
19     def __init__(self, input_size, hidden_size, output_size, num_layers):
20
21         super(Decoder, self).__init__()
22
23         self.lstm = torch.nn.LSTM(input_size = input_size, hidden_size = hidden_size,
24                                   num_layers = num_layers,
25                                   batch_first = True)
26
27         self.fc_decoder = torch.nn.Linear(hidden_size, output_size)
28
29     def forward(self, input_seq, encoder_hidden_states):
30
31         lstm_out, hidden = self.lstm(input_seq, encoder_hidden_states)
32         output = self.fc_decoder(lstm_out)
33
34         return output, hidden
```



# Define Model



Combine Encoder and Decoder classes into a single class (Encoder\_Decoder)

```
36 class Encoder_Decoder(torch.nn.Module):
37
38     def __init__(self, input_size, hidden_size, decoder_output_size, num_layers):
39
40         super(Encoder_Decoder, self).__init__()
41
42         self.Encoder = Encoder(input_size = input_size, hidden_size = hidden_size,
43                                num_layers = num_layers)
44
45         self.Decoder = Decoder(input_size = input_size, hidden_size = hidden_size,
46                                output_size = decoder_output_size, num_layers = num_layers)
```



# Define Hyperparameters

```
1 Encoder_Decoder_RNN = Encoder_Decoder(input_size = 1, hidden_size = 15,  
2                                     decoder_output_size = 1, num_layers = 1)  
3  
4 learning_rate = 0.01  
5 epochs = 50  
6  
7 batchsize = 5  
8 num_features = train_output_seqs.shape[2]  
9  
10 loss_func = torch.nn.MSELoss()  
11 optimizer = torch.optim.Adam(Encoder_Decoder_RNN.parameters(), lr=learning_rate)  
12  
13 Encoder_Decoder_RNN
```

Define Encoder Decoder Specifics

Define Learning rate, epochs,  
batchsize and num\_features/timestep

Define loss function and optimizer



# Identify Tracked Values

```
1 train_loss_list = []
```

Empty Python list for keeping track of loss values



# Train Model

```
1 train_input_seqs = torch.from_numpy(train_input_seqs).float()
2 train_output_seqs = torch.from_numpy(train_output_seqs).float()
3
4 train_batches_features = torch.split(train_input_seqs, batchsize)[: -1]
5 train_batches_targets = torch.split(train_output_seqs, batchsize)[: -1]
6
7 batch_split_num = len(train_batches_features)
```

Convert numpy arrays to torch tensors

Split training data into mini-batches  
(skip last mini-batch since it can have  
smaller batch size)

Compute total number of mini-batches



# Train Model

```
1 for epoch in range(epochs): # For each epoch
2
3     for k in range(batch_split_num):
4
5         hidden_state = None
6
7         decoder_output_seq = torch.zeros(batchsize, decoder_outputseq_len, num_features)
8
9         optimizer.zero_grad()
10
11         encoder_output, encoder_hidden = Encoder_Decoder_RNN.Encoder(train_batches_features[k], hidden_state)
12         decoder_hidden = encoder_hidden
13
14         decoder_input = train_batches_features[k][:, -1, :]
15         decoder_input = torch.unsqueeze(decoder_input, 2)
16
17         for t in range(decoder_outputseq_len):
18
19             decoder_output, decoder_hidden = Encoder_Decoder_RNN.Decoder(decoder_input, decoder_hidden)
20
21             decoder_output_seq[:, t, :] = torch.squeeze(decoder_output, 2)
22
23             decoder_input = train_batches_targets[k][:, t, :]
24             decoder_input = torch.unsqueeze(decoder_input, 2)
25
26         loss = loss_func(torch.squeeze(decoder_output_seq), torch.squeeze(train_batches_targets[k]))
27
28         train_loss_list.append(loss.item())
29
30         loss.backward()
31
32         optimizer.step()
33
34     print("Averaged Training Loss for Epoch ", epoch, ": ", np.mean(train_loss_list[-batch_split_num:]))
```

Define initial hidden states and empty tensor for decoder outputs

Pass training input sequence + hidden states to encoder

Initial input to decoder = last value of the input sequence

Fill in decoder output tensor by using teacher forcing method (provide ground truth inputs)

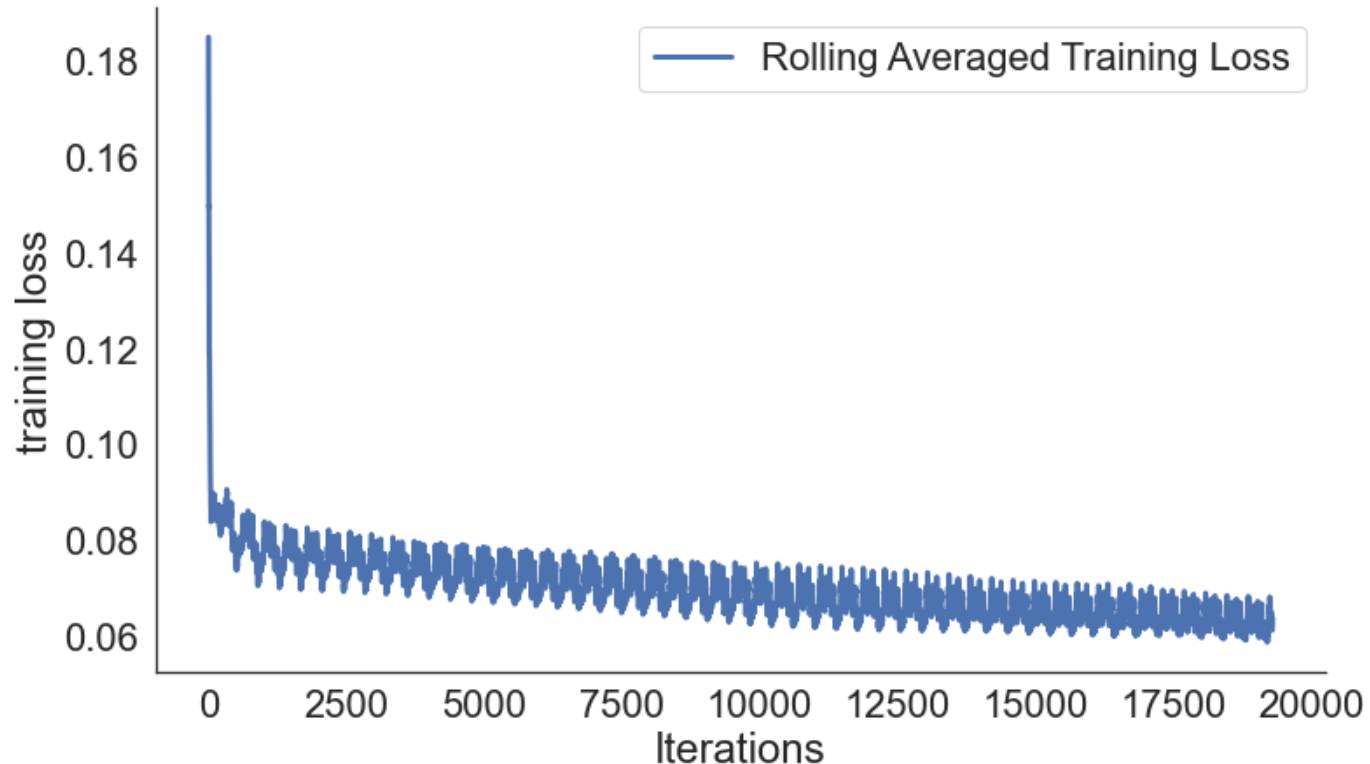
Compute and append Loss  
Back-propagation  
Update network



# Visualize & Evaluate Model


```
1 plt.figure(figsize = (12, 7))
2
3 plt.plot(np.convolve(train_loss_list, np.ones(100), 'valid') / 100,
4          linewidth = 3, label = 'Rolling Averaged Training Loss')
5 plt.ylabel("training loss")
6 plt.xlabel("Iterations")
7 plt.legend()
8 sns.despine()
```

Plot moving average training loss





# Visualize & Evaluate Model



See example notebook for detailed code implementation





# LAB 6 ASSIGNMENT:

Stock Prediction AI with Encoder-Decoder RNN



# Stock Dataset

	Date	Open	High	Low	Close	Adj Close	Volume
0	2010-06-29	19.000000	25.00	17.540001	23.889999	23.889999	18766300
1	2010-06-30	25.790001	30.42	23.299999	23.830000	23.830000	17187100
2	2010-07-01	25.000000	25.92	20.270000	21.959999	21.959999	8218800
3	2010-07-02	23.000000	23.10	18.709999	19.200001	19.200001	5139800
4	2010-07-06	20.000000	20.00	15.830000	16.110001	16.110001	6866900

	Date	Open	High	Low	Close	Adj Close	Volume
0	2004-08-19	50.050049	52.082081	48.028027	50.220219	50.220219	44659000
1	2004-08-20	50.555557	54.594593	50.300301	54.209209	54.209209	22834300
2	2004-08-23	55.430431	56.796795	54.579578	54.754753	54.754753	18256100
3	2004-08-24	55.675674	55.855854	51.836838	52.487488	52.487488	15247300
4	2004-08-25	52.532532	54.054054	51.991993	53.053055	53.053055	9188600

	Date	Open	High	Low	Close	Adj Close	Volume
0	1985-01-29	1277.719971	1295.489990	1266.890015	1292.619995	1292.619995	13560000
1	1985-01-30	1297.369995	1305.099976	1278.930054	1287.880005	1287.880005	16820000
2	1985-01-31	1283.239990	1293.400024	1272.640015	1286.770020	1286.770020	14070000
3	1985-02-01	1276.939941	1286.109985	1269.770020	1277.719971	1277.719971	10980000
4	1985-02-04	1272.079956	1294.939941	1268.989990	1290.079956	1290.079956	11630000


- TSLA.csv
- 2227 days
- 7 attributes

- GOOGL.csv
- 3702 days
- 7 attributes

- DJI.csv
- 8636 days
- 7 attributes



# Stock Prediction AI with Encoder-Decoder RNN



In this exercise, you will use Encoder-Decoder RNN architecture to predict the **last 100 days'** stock values.

You are free to pick **one of the stock datasets (TSLA, GOOGL, DJI)** for training and testing your model. Use **closing stock value (i.e., "Close" column)** for both training and testing data.

Feel free to pick **encoder/decoder sequence sizes** of your choice, **LSTM** or **GRU** for your RNN cell as well as **RNN extensions** such as Deep RNN or Bi-directional RNN.

Before training, **normalize the data** and create **train\_input\_seqs** and **train\_output\_seqs** like the example task.

After training, plot your **RNN predicted stock value** against the **ground truth test values** and calculate its **MSE error**. Use **Teacher forcing method** for predicting test outputs.